EXPLORING THE USE OF THE SEMANTIC WEB FOR DISCOVERING AND PROCESSING DATA FROM SENSOR OBSERVATION SERVICES

A thesis submitted to the Delft University of Technology in partial fulfillment
of the requirements for the degree of

Master of Science in Geomatics

by

Ivo de Liefde

June 2016

The work in this thesis was made in the:

Geo-Database Management Centre
Department of the OTB
Faculty of Architecture & the Built Environment
Delft University of Technology

Supervisors: M. de Vries
               B.M. Meijers
Co-reader: S. Zlatanova

ABSTRACT

[Should fit on one page.]

Bacon ipsum dolour sit amet porchetta beef turkey, bacon turducken boudin hamburger venison ball tip. Brisket pork loin bresaola short loin ground round leberkas pastrami tongue jerky cow turducken beef ribs. Pork ribeye landjaeger prosciutto pig venison tenderloin. Swine beef ribs kielbasa, porchetta tenderloin salami venison pork belly tail. Bacon ipsum dolour sit amet porchetta beef turkey, bacon turducken boudin hamburger venison ball tip. Brisket pork loin bresaola short loin ground round leberkas pastrami tongue jerky cow turducken beef ribs. Pork ribeye landjaeger prosciutto pig venison tenderloin. Swine beef ribs kielbasa, porchetta tenderloin salami venison pork belly tail. Bacon ipsum dolour sit amet porchetta beef turkey, bacon turducken boudin hamburger venison ball tip. Brisket pork loin bresaola short loin ground round leberkas pastrami tongue jerky cow turducken beef ribs. Pork ribeye landjaeger prosciutto pig venison tenderloin. Swine beef ribs kielbasa, porchetta tenderloin salami venison pork belly tail.

ACKNOWLEDGEMENTS

Thanks to everyone, especially to my supervisors and my mum. And obviously to the ones who made that great template.

Bacon ipsum dolour sit amet porchetta beef turkey, bacon turducken boudin hamburger venison ball tip. Brisket pork loin bresaola short loin ground round leberkas pastrami tongue jerky cow turducken beef ribs. Pork ribeye landjaeger prosciutto pig venison tenderloin. Swine beef ribs kielbasa, porchetta tenderloin salami venison pork belly tail.

Bacon ipsum dolour sit amet porchetta beef turkey, bacon turducken boudin hamburger venison ball tip. Brisket pork loin bresaola short loin ground round leberkas pastrami tongue jerky cow turducken beef ribs. Pork ribeye landjaeger prosciutto pig venison tenderloin. Swine beef ribs kielbasa, porchetta tenderloin salami venison pork belly tail.

Bacon ipsum dolour sit amet porchetta beef turkey, bacon turducken boudin hamburger venison ball tip. Brisket pork loin bresaola short loin ground round leberkas pastrami tongue jerky cow turducken beef ribs. Pork ribeye landjaeger prosciutto pig venison tenderloin. Swine beef ribs kielbasa, porchetta tenderloin salami venison pork belly tail.

...

CONTENTS

ACRONYMS

1 | INTRODUCTION

From 2020 onwards all member states of the European Union (EU) should provide sensor data to the Infrastructure for Spatial Information in Europe (INSPIRE) in order to comply with annex II and III of the INSPIRE directive [INSPIRE, 2015]. For this a number of Sensor Web Enablement (SWE) standards are required to be used [INSPIRE, 2014]. The sensor web is a relatively new development and there are still many questions on how to structure it. This thesis aims to design a method to publish and link sensor metadata on the semantic web to improve the discovery, integration and aggregation of sensor data using SWE standards.

1.1 BACKGROUND

In 2008 the Open Geospatial Consortium (OGC) introduced a new set of standards called Sensor Web Enablement (SWE). These standards make it possible to connect sensors to the internet and retrieve data in a uniform way. This allows users or applications to retrieve sensor data through standard protocols, regardless of the type of observations or the sensor's manufacturer [Botts et al., 2008]. Among other standards SWE includes the Observations and Measurements (O&M) which is a model for encoding sensor data, the Sensor Modelling Language (SensorML) which is a model for describing sensor metadata and the Sensor Observation Service (SOS) which is a service for retrieving sensor data [Botts et al., 2007]. O&M has also been adopted by the International Organisation for Standardisation (ISO) under ISO 19156:2011 [ISO, 2011].

Recently OGC has defined the role which their standards could play in smart city developments [Percivall, 2015]. Smart cities can be defined as "enhanced city systems which use data and technology to achieve integrated management and interoperability" [Moir et al., 2014, p. 18]. Research on smart cities has shown a great potential for using sensor data in urban areas. Often this is presented in the context of the Internet of Things (IoT) [Zanella et al., 2014; Wang et al., 2015a]. The IoT can be described as "the pervasive presence around us of a variety of *things* or *objects* ... [which] are able to interact with each other and cooperate with their neighbors to reach common goals" [Atzori et al., 2010, p. 2787].

Parallel to the development of the sensor web other research has focused on the semantic web, as proposed by Berners-Lee et al. [2001]. This is a response to the traditional way of using the web, where information is only available for humans to read. The semantic web is an extension of the internet which contains meaningful data that machines can understand as well. Rather than publishing documents on the internet the semantic web contains linked data using the Resource Description Framework (RDF), also known as the *web of data* [Bizer et al., 2009]. Data in RDF can be queried using the SPARQL Protocol and RDF Query Language (SPARQL) at so-called

SPARQL endpoints. The Web Ontology Language (OWL) is an extension of RDF and was designed "to represent rich and complex knowledge about things, groups of things, and relations between things" [OWL working group, 2012]. Originally, the semantic web intended to add metadata to the internet [Lassila and Swick, 1999]. However, today it is being used for linking any kind of data from one source to another in a meaningful way [Cambridge Semantics, 2015].

Sheth et al. [2008] proposes to use semantic web technologies in the sensor web. This Semantic Sensor Web (SSW) builds on standards by OGC and the World Wide Web Consortium (W3C) "to provide enhanced descriptions and meaning to sensor data" [Sheth et al., 2008, p. 78]. W3C responded to this development by creating a standard ontology for sensor data on the semantic web [Compton et al., 2012].

1.2 PROBLEM STATEMENT

Finding sensor data that can be retrieved using open standards is not easy. The implementation of the sensor web is still in an early stage. At the moment there are only a limited number of SOS implementations available on the web and they contain a limited amount of data. In the Netherlands the SOS by the Dutch national institute for public health and the environment (RIVM) is one of the first ones to be developed. It has only recently been launched and contains data on air quality. A number of other organisations still use a custom Application Programming Interface (API) to retrieve data from sensors connected to the internet. The problem of these custom APIs is that it is very hard to create an application that automatically retrieves data from them, because they have not implemented standards regarding the content of their service, the metadata models behind it or the kind of requests that can be made. It forces the application to have knowledge built in on the specifics of the individual APIs that are being used.

It has been researched to what extent a catalogue service could be useful for discovering sensor data from a SOS using the web service interfaces Sensor Instance Registry (SIR) [Jirka and Nüst, 2010] and Sensor Observable Registry (SOR) [Jirka and Bröring, 2009]. Catalogue services have already been available for example for the Web Map Service (WMS), Web Feature Service (WFS) or Web Coverage Service (WCS) [Nebert et al., 2007]. However, for the sensor data sources used in this paper no register or catalogue service has been implemented. Atkinson et al. [2015] also argues that catalogue services have a number of major disadvantages. It places a very high burden on the client to not only know where to find the catalogue service, but also to have knowledge on all kinds of other aspects (e.g. its organisation, access protocol, response format and response content) [Atkinson et al., 2015, p. 128]. Atkinson et al. suggest that linked data is therefore a much better solution for discovering sensor data.

However, for sensor data to be discovered on the semantic web there have to be inward links, from other sources linking towards the sensor (meta)data. Current research on the SSW has focused on publishing sensor data on the semantic web with links that point outwards [Atkinson et al., 2015; Janowicz et al., 2013; Pschorr, 2013]. This gives meaning to the data and is useful in order to work with the data, but it has a very limited effect on the discovery of the sensor data by others.

One of the challenges of using sensor data is the difficulty of integrating it from different sources to perform data fusion [Corcho and Garcia-Castro, 2010; Ji et al., 2014; Wang et al., 2015b]. Data fusion is "a data processing technique that associates, combines, aggregates, and integrates data from different sources" [Wang et al., 2015a, p. 2]. Even if the sources comply with the SWE standards it is challenging, since the data can be of a different granularity, both in time and space. Spatio-temporal irregularities are a fundamental property of sensor data [Ganesan et al., 2004].

The question arises to what extent the semantic web could be a better solution for publishing sensor data than the current geoweb solutions like SOS. The geoweb has some very good qualities, such as very structured approaches through which (sensor) data can be retrieved using well defined services. These standardised services have been accepted by large organisations as OGC and ISO. Furthermore, they are often based on years of discussion. This is different from for example web pages where content can be completely unstructured. The response of a SOS also contains some semantics about sensor data. There can be x-links inside the Extensible Markup Language (XML) with Uniform Resource Identifier (URI)s that point to semantic definitions of objects.

Still, the semantic web could be beneficial for the geoweb. Since data on the web has a distributed nature it can be questioned whether centralised catalogue services are feasible to create. It places a burden on the owner of the SOS to register with a catalogue service. Also, there could be multiple of these services on the web creating issue regarding the discovery of relevant catalogues. The semantic web could solve this issue by getting rid of the information silos and storing data directly on the web instead. This allows the interlinking and reuse of data on the web, which makes it easier to find related data. For automatic integration and aggregation it could be useful that the semantic web is machine understandable.

In conclusion, the problem to be addressed is the lack of knowledge on how to exploit the full potential of the sensor web using the semantic web. Creating the right links could greatly enhance the discovery, integration and aggregation of sensor data. However, there is no method yet to establish this linked metadata for sensors, while the standardised nature of a SOS should allow for generating it in an automated process. This thesis will create a design for such an automated process, research how to establish inward links and explore the advantages and disadvantages of publishing sensor metadata on the semantic web with a proof-of-concept implementation.

1.3 MOTIVATION

Sensor data ties together many different fields of research. On the one hand there is research on how to create the most efficient sensor networks that uses the least amount of power to transfer the observed data over long distances [Korteweg et al., 2007; Xiang et al., 2013]. This involves academic fields such as mathematics, physics and electrical engineering. On the other hand there is research that uses sensor data to gain insights into real world phenomenon. This involves academic fields such as geography, environmental studies and urbanism. In order to connect these scientific fields, studies have focused on the use of computer science and standardisation for transferring sensor data over the internet.

In the future more sensor data is expected to be produced [Price Waterhouse Coopers, 2014]. Both experts and non-experts will be involved in this development. Experts will produce more data because of European legislation (INSPIRE). Non-experts will be involved more often via smart cities and IoT developments where users or consumer electronics produce sensor data as well. This vast amount of data could be very useful for academic research, provided researchers are able to find the data they need online and are able to integrate and aggregate data from heterogeneous sources. Publishing sensor metadata on the semantic web could make it easier to find what you need through related data on the internet. Having a automated process for this and being able to seamlessly integrate and aggregate data from different sources could be of great use for research such as van der Hoeven et al. [2014], Van der Hoeven and Wandl [2015] and Theunisse [2015]. They are examples of studies that try to understand phenomenon in the built environment using sensor data. Currently data collection and processing takes up a large part of the research, while with the implementation of SWE standards and the use of the semantic web this might be significantly reduced.

Sensor data is becoming more important in the every day life of many people. Therefore, the gap between the logical data requests of users and the technical requests defined by sensor web protocols should be bridged. A logical sensor data request contains at least a geographical area, a type of observation and a time range. An example of a simple logical request is: The average air temperature per month over the last five years of sensors located in Delft. To translate this to a technical request specific knowledge is required about things like URIs, encodings, data models, service models and data formats. This thesis looks into the use of the semantic web to allow sensor data to be easily retrieved from the internet, by understanding the user's logical request and executing it automatically.

1.4 RESEARCH QUESTIONS

This thesis aims to design a method that uses the semantic web to improve sensor data discovery as well as the integration and aggregation of sensor data from multiple sources. The following question will be answered in this research:

> *To what extent can the semantic web improve the discovery, integration and aggregation of distributed sensor data?*

Subquestions:

- To what extent can sensor metadata be automatically retrieved from any SOS?

- To what extent can sensor metadata from a SOS be automatically converted to linked data and published on the semantic web?

- What is an effective balance between the semantic web and the geo web in the chain of discovering, retrieving and processing sensor data?

- To what extent can already existing standards for retrieving data be (re)used for a service that supplies integrated and aggregated sensor data?

2 | RELATED WORK

A number of research topics are relevant for this thesis: how to use existing standards for publishing sensor data to the semantic web, developing ontologies that are suitable for many different kinds of sensor data and how to aggregate sensor data based on geographical features and time. This chapter discusses the recent relevant literature on these topics. However, first it will introduce the Sensor Web Enablement suite of standards and the semantic web.

2.1 SENSOR WEB ENABLEMENT

Botts et al. [2007] present Sensor Web Enablement (SWE), which is a suite of standards developed by OGC. It contains two parts: the information model and the service model. The information model includes O&M, SensorML and SWE common. O&M defines the data model and encoding for observation data and SensorML defines the data model and encoding for sensor metadata. SWE common is a low-level data model for exchanging sensor related data. The service model contains the SOS, Sensor Planning Service (SPS), Sensor Alert Service (SAS) and Web Notification Service (WNS). A SOS can be used to retrieve observation data, a SPS to plan actions of a sensor, and a SAS to receive alerts about subscribed events. With a WNS users can have asynchronous dialogues (message interchanges) with one or more other services. Recently the SensorThings API has been added to the list of SWE services (see http://ogc-iot.github.io/ogc-iot-api/index.html). The Sensor-Things API is a service for retrieving observation data and sensor metadata for IoT applications. This thesis focusses on O&M, SensorML and SOS. The following paragraphs will therefore described these standards in more detail.

2.1.1 Observation and Measurements

In the Observations and Measurements (O&M) data model an observation is modelled using a number of concepts. First of all, there is a feature of which sensor data is wanted. This is the so-called Feature of Interest (FOI). This feature has a property that can be observed, also known as an osbervable property. For example, the air at a certain location (FOI) has a certain temperature (observable property). A FOI can also be a geographical feature such as a river or a forest area. O&M uses the concept of FOI rather than the sensor location, because for some observations the exact location may not be trivially available. Furthermore, specific specimens are in some cases removed from their sampling location and observed at another location afterwards. This would create problems when defining a location of observation, hence the use of Feature of Interests in the O&M data model. An observation is defined as the "act of measuring or otherwise determining the value of a property" [ISO, 2011, p. 3]. The observation uses a procedure, which is a
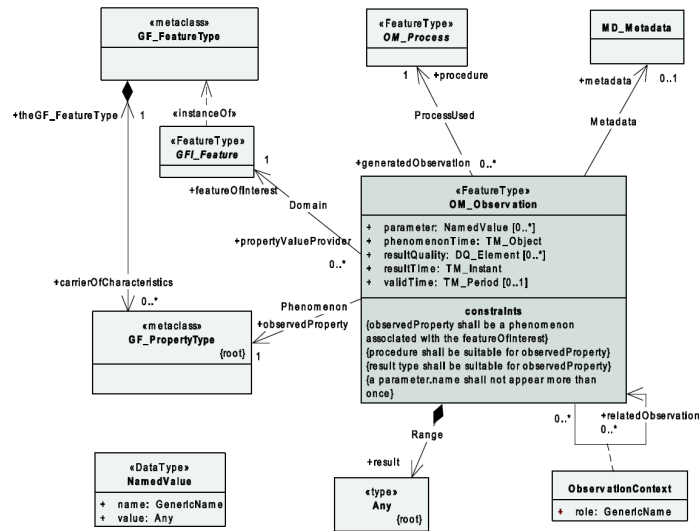
Figure 2.1: Basic observation in the O&M data model [ISO, 2011, p. 9]

method, algorithm or instrument, or system of these. A procedure can therefore be a sensor, an algorithm processing the raw observation data and/or a system of sensors observing a property of the FOI. The output of the procedure is a result. The result consists of a value and a corresponding Unit of Measurement (UOM). For example: 15 (value) degrees Celcius (UOM).

A procedure can produce many results about the same FOI over time. Therefore, three temporal concepts have been modelled: result time, phenomenon time and valid time. The result time is the timestamp of the observation result, or in other words: the moment the observation result became available. It is not necessarily the time that corresponds to the phenomenon, depending on when the observation procedure is performed (specimens can be taken from a sampling location and observed later) and the amount of time it takes to perform a procedure. Therefore, the phenomenon time describes the time that the observation applies to the property of the FOI. The valid time is a time range that indicates when an observation is a valid indication of the property of the FOI. When observing glacier motion at a speed of several meters per year an observation result might be valid for a day. On the other hand, observations with large fluctuations during the day might only be valid for a number of minutes up to an hour. This is the case for example with air quality observations in cities, that have large fluctuations with peaks during the rush hours.

Additional metadata about an observation can be added using the metadata class in the data model. To provide context about an observation the 'related observations' class has been added to the O&M model as well. Figure 2.1 shows the Unified Modeling Language (UML) diagram containing the different concepts explained in this paragraph and the relations between them.

2.1.2 SensorML

The Sensor Modelling Language (SensorML) aims to "provide a robust and semantically-tied means of defining processes and processing components associated with the measurement and post-measurement transformation of observations" [OGC, 2014, p. ix]. Creating interoperability at the syntactic and semantic level allows processes to be better understood by machines,

Figure 2.2: 'DescribedObject' class in SensorML [OGC, 2014, p. 39]

utilized automatically in complex workflows, and to be easily shared between intelligent sensor web nodes. SensorML is a framework that can be used to define the geometric, dynamic, and observational characteristics of sensors and sensor systems in order to achieve syntactic and semantic interoperability.

In SensorML there is a destinction between two types of processes: physical and non-physical processes. For a physical process information about the spatio-temporal position is important. This is the case with for example detectors, actuators, and sensor systems. Non-physical processes are mathematical operations or functions. A process is modelled as a specialisation of the 'DescribedObject' class. This class provides a set of metadata which are useful for all process classes in SensorML. Figure 2.2 shows the properties of a described object class instance. The 'AbstractProcess' class is derived from the described object and includes the properties: inputs, outputs, parameters, typeOf, featureOfInterest, configuration and modes. This forms the basis for defining both physical and non-physical processes. The 'AbstractPhysicalProcess' class is derived from the describe object class and adds spatial and temporal coordinates for the physical process device. The final step in defining a physical process is to specialise the abstract physical process with the 'PhysicalComponent' class. The physical component describes the device that provides a processing function. Figure 2.3 shows the UML class diagram of a physical process.

Although SensorML is not dependent on O&M the result of a process modelled by SensorML is typically considered as an observation result. This is the case if it is measuring a physical property or phenomenon. Therefore, the output values described in SensorML and resulting from a sensor or process may be encoded as an O&M Observation. Inversely, the procedure property of an O&M Observation instance can be described using SensorML [OGC, 2014].

2.1.3 Sensor Observation Service

There are three core requests that can be made to retrieve sensor (meta)data from a SOS: `GetCapabilities`, `DescribeSensor` and `GetObservation`. `Get Capabilities` returns a complete overview of what the SOS has to offer. The `DescribeSensor` request returns detailed information about individ-
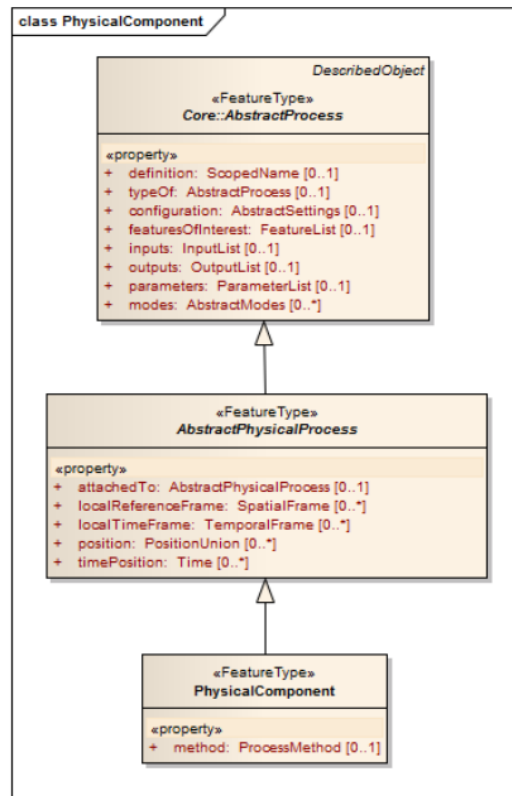
Figure 2.3: Definition of a physical process in SensorML [OGC, 2014, p. 57]

ual sensors. These three core requests are mandatory in a SOS under the 2.0 specifications [Bröring et al., 2012]. There are also a number of optional extensions to a SOS: `GetFeatureOfInterest`, `GetObservationById`, `InsertCapabilities`, `InsertObservation`, `InsertSensor`, `DeleteSensor`, `InsertResult`, `InsertResultTemplate` and `GetResultTemplate`. Requests can be made as a HyperText Transfer Protocol (HTTP) GET request or a HTTP POST request. There can be different response formats, but there is always at least the option to retrieve the response as an XML document. Based on the specification by Bröring et al. [2012] the following paragraphs describe the core and optional requests of a SOS, as well as the structure of their responses.

Get capabilities

The `GetCapabilities` request is the first step in communicating with a SOS. The request is made by taking the HTTP address of the SOS and adding `service=SOS&request=GetCapabilities`. It returns a document including information on what the service has to offer. The document contains a number of sections: service identification, service provider, operations metadata, filter capabilities and contents.

In the service identification section there is general information about the service, such as the title and supported SOS versions, but also whether there are fees or access constraints. The service provider section contains details on which organisation provides the SOS and lists their contact information. The operations metadata section lists the supported request types. It also contains an overview of all features-of-interest, observed properties, pro-

Figure 2.4: Data model behind the capabilities document of a SOS [Bröring et al., 2012]

cedures and offerings. Offerings are similar to layers in a WMS, grouping together observations collected by one procedure.

The contents section describes the data that can be retrieved, grouped in offerings. Each offering has an identifier, together with information on the procedure, observable properties and the feature-of-interest type. Which filters can be applied in a request is described in the filter capabilities section. The supported parameters for both spatial and temporal filters are listed here.

Describe sensor

The `DescribeSensor` request gives detailed information on a specific sensor. The request is built by taking the HTTP address of the SOS and adding `service=SOS&version=2.0.0&request=DescribeSensor& procedure=aprocedure&proceduredescriptionformat=aformat` where the procedure and procedure description format have to contain values defined in the capabilities document.

Get observation

Using `GetObservation` actual measurements can be retrieved. The request is made by taking the HTTP address of the SOS and adding `service=SOS& version=2.0.0&request=GetObservation`. This returns a response with the default parameters, which can differ from one SOS to another. To further specify the request, paramaters can be added such as: observed property, procedure, feature-of-interest, offering and outputformat. Spatial and temporal filters can be added if these are supported by the service.

The `GetObservationByID` request is an extension that let's users retrieve an observation using an identifier that points to a certain observation.

Get feature of interest

The `GetFeatureOfInterest` request allows to retrieve information about the feature of interest of a certain observation. The response can be all features of interest or only ones that are related to a specific observed property, procedure or spatial filter. This request also allows logical operators, for example: "GetFeatureOfInterest (observedProperty := temperature AND procedure := thermometerX OR anemometerY)" [Bröring et al., 2012, p. 40]. The response is a document with 'GFI_features' (ISO 19109) [International Organisation for Standardisation, 2005], which are implemented in the Geography Markup Language (GML) (ISO 19136) [International Organisation for Standardisation, 2007] by the element gml:AbstractFeature and type gml:AbstractFeatureType [ISO, 2011, p. 38].

Transactional extensions

It is possible to update the content of a SOS using transactional request. There are six of these requests that could be implemented: `InsertCapabilities`, `InsertObservation`, `InsertSensor`, `DeleteSensor`, `InsertResultTemplate` and `InsertResult`. `InsertCapabilities` allows a request to add data to the capabilities document described in Paragraph 2.1.3. The request contains three mandatory parameters and one optional parameter: it is required to have the `procedureDescriptionFormat`, `FeatureOfInterestType` and `ObservationType`. Optionally, `SupportedEncoding` could be added to the request. An `InsertCapabilities` should be made in combination with a `InsertObservation`, `InsertSensor` or `InsertResult` request.

A `InsertResultTemplate` request allows to upload a template for result values. It should contain data about the offering, the observation template, the result structure and result encoding. The actual results can be added later using an `InsertResult` request. This request is different from the `InsertObservation` request, as it only inserts the result value of an observation, assuming that the metadata is already present in the SOS. This is useful when there is limited communication bandwidth and processing power. This request has two mandatory parameters: a pointer to the template and the observation value to be inserted. An `InsertObservation` request allows observations to be added to a registered sensor system. It also has two mandatory parameters: a pointer to an offering and the observation to be inserted.

Individual sensors can be inserted or deleted using `InsertSensor` and `DeleteSensor` requests. For inserting a sensor the following parameters are required: the procedure description format, a procedure description, the observable property, a feature of interest type and an observation type. For deleting a sensor an identifier that points to a specific sensor needs to be passed as a parameter.

2.2 SEMANTIC WEB

2.2.1 Resource Description Framework

In RDF data is stored as so-called 'triples'. These triples are structured as: subject, predicate and object [Berners-Lee et al., 2001]. The subject and the object are things and the predicate is the relation between these two things.
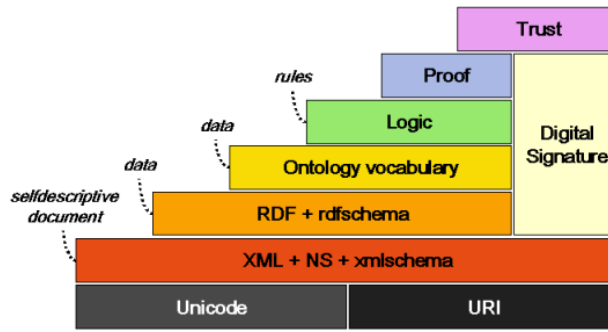
Figure 2.5: Hierarchy of the semantic web [Koivunen and Miller, 2002]

Figure 2.6: Triples of object, predicate and subject define Delft as a municipality with a geometry

For example, to define a geographic feature such as the municipality of Delft on the semantic web a number of triples can be made. Figure 2.6 shows how Delft can be defined as a municipality with a certain geometry using triples of subject, predicate and object.

Three types of data can make up these triples [Manola et al., 2014]. The first type is an International Resource Identifier (IRI). This is a reference to a resource and can be used for all positions of the triple. A Uniform Resource Locator (URL) is an example of an IRI, but IRIs can also refer to resources without stating a location or how it can be accessed. An IRI is a generalisation of an URI, and also allows non-ASCII characters. In the example of the municipality of Delft, IRIs can be used to define 'Delft' and 'Municipality', but also for the predicates 'is a' and 'has geometry'. The second type of data is a literal. A literal is a value which is not an IRI, such as strings, numbers or dates. These values can only be used as object in a triple. In the example of Delft, a literal could be used to store the actual geometry of the boundary: POLYGON((x_1, y_1 x_2, y_2 ... x_n, y_n, x_1, y_1)). A literal value can have a datatype specification [Cyganiak et al., 2014]. This is added to the literal with the ^^ symbols, followed by the IRI of the datatype specification. In Figure 2.7 the datatype is 'geo:wktLiteral'.

Sometimes it is useful to refer to things without assigning them with a global identifier. The third type is the blank node and can be used as an subject or object without using an IRI or literal [Manola et al., 2014].

2.2.2   Notation

There are a number of different notations for writing down these triples (serialisation), such as XML [Gandon and Schreiber, 2014], N3 [Berners-Lee and Connolly, 2011] and Turtle [Beckett et al., 2014]. Turtle will be used in this thesis, because it is commonly used notation which is also relatively easy to read for humans. The DBPedia IRI is used for the object 'Municipality'. The 'is a' predicate is represented by a built-in RDF predicate which can

```
PREFIX geo: <http://www.opengis.net/ont/geosparql#>

<http://example.com/Delft> a <http://dbpedia/resource/Municipality> ;
    geo:hasGeometry "<http://www.opengis.net/def/crs/EPSG/0/4258> POLYGON(( x1 y1, x2 y2, ... xn yn, x1 y1 ))"^^geo:wktLiteral
```

Figure 2.7: Triples of Figure 2.6 in the Turtle notation

Figure 2.8: Persistent Uniform Resource Locator (PURL) resolves to the current resource location [Shafer et al., 2016]

be written simple as 'a'. The second predicate is 'hasGeometry' for which the GeoSPARQL IRI is used. The geometry is a literal in the Well-Known Text (WKT) format. Note that the subject is only written once when there are multiple triples with the same subject. Triples that share the same subject are divided by semicolons. A point marks the end of the last triple with a specific subject.

2.2.3 Persistent Uniform Resource Locators

URLs are an essential part of the web. However, if an URL changes the existing links towards this URL are broken. To prevent this Persistent Uniform Resource Locators (PURLs) are being used. A Persistent Uniform Resource Locator is a "naming and resolution service for general Internet resources" [Shafer et al., 2016]. This allows organisations to change the location of their data without changing the URL to which can be linked. A PURL server receives the URL and redirects the client to the current location of the resource. If the location of the resource changes, the server can be informed. It will then redirect clients to the new location (Figure 2.8).

2.3 SENSOR METADATA IN CATALOGUE SERVICES

The OGC has developed a way of including sensors into a catalogue service using sensor registries. The different services related to this are described in this paragraph, starting with the Catalog Service for the Web (CSW), after which the Sensor Observable Registry (SOR) and Sensor Instance Registry (SIR) services will be presented.

2.3.1 Catalog Service for the Web

The CSW is an OGC standard for a geoweb service that contains "collections of descriptive information (metadata) for data, services, and related information objects" [Open Geospatial Consortium, 2007, p. xiv]. Figure 2.9 shows
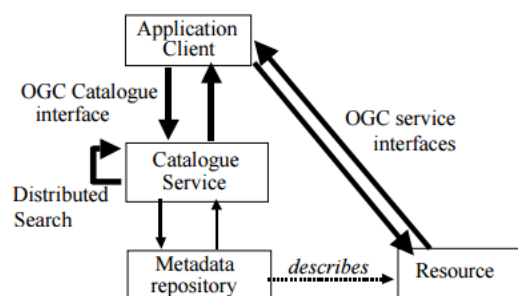
Figure 2.9: CSW model architecture [Open Geospatial Consortium, 2007, p. 26]

the intermediary role of this service between client and data sources. It also shows that the CSW can use data from three different sources to return to a client: a local metadata repository, a resource service, or another CSW. The resource service can use an OGC interface, but this is not required.

The CSW has the following requests for retrieving metadata: GetCapabilities, DescribeRecord, GetDomain, GetRecords and GetRecordByID. Metadata producers can use the Transaction and Harvest requests to respectively change or add content of a CSW. The GetCapabilities request returns a capabilities document to the client showing what the service has to offer. This request and response is required for all OGC geoweb services and contains sections such as: service identification, service provider, operations metadata and filter capabilities. The service identification section lists general information about the service, such as the title and supported CSW versions, but also whether there are fees or access constraints. The service provider section contains details on which organisation provides the SOS and lists their contact information. The operations metadata section lists the supported request types. The filter capabilities show the different filters that are supported by the CSW instance.

The DescribeRecord request allows the client to recieve a description of (a part of) the information model inside the catalog service. Parameters for namespaces or type names can be added to retrieve a part of the information model. The optional GetDomain request can be used to retrieve information about the range of values for a metadata record element or request parameter. A GetRecords request can be made to search for and retrieve catalogue records or to see if a they are present. The query element is the encoding for the search part of the GetRecords. The constraint parameter contains the query element and the constrain language parameter set the language that is being used to make the query. To see whether a record is present the outputSchema parameter and the 'ElementName' or 'ElementSetName' parameter(s) should be used. The GetRecordsByID request returns records using their identifier.

The Transaction request can be used to create, modify and delete catalogue records. HTTP GET requests using Key-Value Pair (KVP) are not supported for this operation, because it is not a convenient way to encode the transaction payloads. Instead POST requests should be used exclusively for this. The transaction element in the request contains the type of action: insert, update and/or delete.

Figure 2.10: Overview of the interaction between CSW, SIR, SOR and potential users [52 North, 2016]

2.3.2 Sensor Observable Registry

The SOR is "a web service interface for managing the definitions of phenomena measured by sensors as well as exploring semantic relationships between these phenomena" [Jirka and Bröring, 2009, p. vi]. This is a web service developed by OGC to enable semantic reasoning on sensor networks, especially concerning phenomenon definitions. This should make it easier to discover sensors that observe a certain phenomenon and to interpret sensor data.

The SOR has four different requests: GetCapabilities, GetDefinitionURNs, GetDefinition and GetMatchingDefinitions. The GetCapabilities request provides an overview of what the SOR has to offer. The capabilities document that is returned to the client contains four sections. These sections are required for every OGC geoweb service and the first three are the same as described in Paragraph 2.3.1. The fourth section is the content section. This part of the capabilities document contains the following information: the number of entries inside the SOR instance, keywords describing the content of the SOR instance, the application domain for which a specific SOR can be applied and an 'ontologyRepositoryURL'. This URL points to a repository that contains an ontology used by this SOR.

GetDefinitionURNs returns a list of Uniform Resource Names (URNs) identifying the definitions that are present in the SOR. Optionally a client can add a 'SearchSensor' parameter to filter URNs. This parameter takes a substring that shall occur within the definition URNs to be returned. Also, a maximum limit for the amount of returned URNs can be added using the 'maxNumberOfResults' parameter. The optional parameter 'startResultElement' can be used to input the number of the first returned result element. For retrieving the definition of a specific URN a GetDefinition request can be made. This request takes an URN of the phenomenon for which a definition should be retrieved as input. The definition is returned as a GML dictionary entry.

GetMatchingDefinitions allows clients to retrieve definitions of observables which in some way are related to another given phenomenon. The re-

lations 'generalization', 'specialization' and 'equivalency' are currently supported for finding matching definitions [Jirka and Bröring, 2009] and have to be specified in the request using the 'matchingtype' parameter. This parameter can take one of three values: SUPER_CLASS, EQUIVALENT_CLASS or SUB_CLASS. Additionally, the request can have the 'searchDepth' parameter, which represents the maximum amount of steps that are allowed in case of a transitively related phenomena.

2.3.3 Sensor Instance Registry

Another web service interface specification by OGC is SIR. SIR is aimed at "managing the metadata and status information of sensors" [Jirka and Nüst, 2010, p. xii]. The goal of this web service is to close the gap between metadata models based on SensorML, which is used in SWE, and the metadata model used in OGC catalogue services. Furthermore, it provides functionalities to discover sensors, to harvest sensor metadata from a SOS, to handle status information about sensors and to link SIR instances to OGC catalogue services.

There are 14 different requests that can be made at a SIR. First of all, there is the GetCapabilities request, which provides an overview of what the service has to offer. For searching and retrieving sensor metadata there are the SearchSensor and DescribeSensor requests. The search sensor request can take a identifier of a specific sensor that is being searched or a search criteria. The criteria can be a value of a metadata property or a spatial filter. The describe sensor request is similar to the one that is supported by SOS (see Paragraph 2.1.3). It returns the SensorML description of a specific sensor.

The HavestService request starts a process to retrieve all available sensor metadata from a specific SWE service. It takes the 'ServiceURL' as input, combined with a service type such as SOS, SPS or SAS. The response document of a harvest request contains a summary of the changes that were performed in the SIR database.

Transactions for individual sensors can be made using InsertSensorInfo, DeleteSensorInfo and UpdateSensorDescription. The insert sensor info request has two mandatory input parameters: 'SensorIDInSIR', which should contain a unique identifier for the inserted sensor and 'SensorDescription', which should contain its metadata. Optionally, a reference to an SWE service can be provided that contains a description of the sensor. The delete sensor info request requires the identifier of the sensor metadata to be deleted, together with a boolean value for whether all data should be deleted or only certain references. In case of the latter the references to be deleted should be provided as well. The update sensor description request contains the identifier of the sensor for which metadata should be updated together with the new sensor description.

There are four requests for managing sensor status information, which are optional for the implementation of a SIR. The GetSensorStatus request is similar to the SearchSensor request, but returns the status of the sensor. To retrieve specific status information a property filter can be added. The SubscribeSensorStatus request allows users to automatically retrieve status information of sensors. It contains the 'SubscriptionTarget' which defines where the status information should be send to. The other parameters specify what kind of status information should be received from which sensors, like a regular GetSensorStatus request. The response includes a subscription identifier and an expiration

date. With the `RenewSensorStatusSubscription` request the expiration date can be extended. It takes the subscription identifier as input. The `CancelSensorStatusSubscription` request allows users to cancel their subscription before the expiration data. Users can also add status information themselves with the `InsertSensorStatus` request.

Managing the connection between a SIR and a OGC catalog can be done using `ConnectToCatalog` and `DisconnectFromCatalog` requests [Jirka and Nüst, 2010]. For connecting to a catalog the URL of the catalog should be provided, combined with time interval. This interval will be used to set the update time: all metadata changes that occurred in the SIR since the previous interval are pushed to the catalog service simultaneously. This connection can be cancelled with a `DisconnectFromCatalog` request.

2.4 SEMANTIC SENSOR DATA MIDDLEWARE

Henson et al. [2009] and Pschorr [2013] suggest adding semantic annotations to a SOS which they call Semantically Enabled SOS (Sem-SOS). In Sem-SOS the raw sensor data goes through a process of semantic annotating before it can be requested with a SOS service. The retrieved data is still an XML document, but with embedded semantic terminology as defined in an ontology. The data retrieved from Sem-SOS is therefore semantically enriched.

Pschorr et al. [2010] has created a prototype that is able to find sensors from a SOS using linked data. The user can input a location and find sensors that are located nearby. They acknowledge the advantages of linked data over a catalogue service. However, the method presented by Pschorr et al. [2010] is still limited to retrieving sensors from a single source in a buffer around a point location.

Janowicz et al. [2013] has specified a method that uses a Representational State Transfer (REST)ful proxy as a façade for SOS. When a specific URI is requested the so-called Semantic Enablement Layer (SEL) translates this to a SOS request, fetches the data and translates the results back to RDF. In this method the sensor data is converted to RDF on-the-fly. This allows the data to be interpreted by both humans and machines.

Atkinson et al. [2015] have identified that "distributed heterogeneous data sources are a necessary reality in the case of widespread phenomena with multiple stakeholder perspectives" [Atkinson et al., 2015, p.129]. Therefore, they propose that methods should be developed to move away from the traditional dataset centric approaches and towards using linked data for cataloguing. This has the potential to bring together data and knowledge from different areas of research about the same (or similar) features-of-interest. It is also argued that using both linked data services and data-specific services could ease the transition into the linked data world.

2.5 SEMANTIC WEB AND THE INTERNET OF THINGS

Barnaghi, Payam and Wang, Wei and Henson, Cory and Taylor, Kerry [2012] describe how the semantic web could be of great importance for the IoT. Jazayeri, Mohammad Ali and Liang, Steve HL and Huang, Chih-Yuan [2015] describe the SOS as one of the protocols that IoT devices could use.

Figure 2.11: The stimulus–sensor–observation pattern [Compton et al., 2012, p. 28]

2.6 SENSOR DATA ONTOLOGIES

Ontologies are necessary to provide meaning to data on the semantic web and to create semantic interoperability. Three recent efforts for developing a standard ontology for sensor data based on SWE standards will be discussed here.

2.6.1 Semantic sensor network ontology

W3C has developed an ontology for sensors and observations called the Semantic Sensor Network Ontology (SSNO). This ontology aims to address semantic interoperability on top of the syntactic operability that the SWE standards provide. To accommodate different definitions of the same concepts the broadest definitions have been used. Depending on the interpretation these can be further defined with subconcepts. The SSNO is based on the stimulus-sensor-observation pattern, describing the relations between a sensor, a stimulus and observations (Figure 2.11). Sensors are defined as "physical objects ... that observe, transforming incoming stimuli ... into another, often digital, representation", stimuli are defined as "changes or states ... in an environment that a sensor can detect and use to measure a property" and observations are defined as "contexts for interpreting incoming stimuli and fixing parameters such as time and location" [Compton et al., 2012, p. 28]. The ontology can be used to model sensor networks from four different perspectives (sensor, observation, system, and feature & property), which they discuss together with additional relevant concepts.

2.6.2 Observation capability metadata model

Hu et al. [2014] have reviewed a number of metadata models (including SensorML and SSNO) for the use of earth observation (including remote sensing). They argue that all of the current metadata models are not sufficient for sensor data discovery. This conclusion is based on an evaluation of six criteria. Three steps were identified in the process of obtaining relevant sensor data for earth observation, which have been used to derive criteria for their evaluation framework. These steps are sensor filtration, sensor optimisation and sensor dispatch. The filtration of sensors should result in a set of sensors that meets the requirements of the application: It should measure

the right phenomenon, be active, be inside the spatial and temporal range, and have a certain sample interval. In sensor optimisation the selected sensors should be combined to complement or enhance each other. To do this, the observation quality, coverage and application is relevant. In the last step – sensor dispatch – the data should be retrieved, stored and transmitted. In every evaluated model the same sensors can be described in different ways or only partially, which affects the outcome of the sensor dispatch.

Therefore, a metadata model is proposed that "reuses and extends the existing sensor observation-related metadata standards" [Hu et al., 2014, p. 10546]. It is composed of five modules: observation breadth, observation depth, observation frequency, observation quality and observation data. They should be derived from metadata elements described using the Dublin Core metadata element set. These five modules can then be formalised following the SensorML schema which can be queried by users via a 'Unified Sensor Capability Description Model-based Engine'.

2.6.3 Om-lite & sam-lite ontologies

Cox [2015b] has been working on new semantic ontologies based on O&M. Previous efforts, such as the SSNO have been using pre-existing ontologies and frameworks. However, there are already many linked data ontologies that could be useful for describing observation metadata, such as space and time concepts. Also, the SSNO does not take sampling features into account. Therefore, Cox [2015b] proposes two new ontologies: OWL for observations or om-lite [Cox, 2015a], which defines the concepts from O&M regarding observations and OWL for sampling features or sam-lite, which defines the sampling feature concepts [Cox, 2015d]. A mapping of the SSNO to om-lite is also provided.

Cox [2015b] describes how the PROV ontology [Lebo et al., 2013] can be directly used inside om-lite. The PROV ontology is "concerned with the production and transformation of Entities through time-bounded Activities, under the influence or control of Agents" [Cox, 2015b, p. 12]. This is a very convenient ontology for modelling real world entities, such as sensors, observation processes and sampling processes. Many other ontologies could be implemented in combination with om-lite and sam-lite, depending on the kind of observations that are being modelled and the data publisher's preference.

2.7 SENSOR DATA AGGREGATION

Sensor data aggregation can be performed for two purposes: To reduce the energy constraint of sensor networks [Korteweg et al., 2007] or to sample a feature-of-interest in space and/or time [INSPIRE, 2014]. Sampling is performed when a feature-of-interest is not accessible, in which case "observations are made on a subset of the complete feature, with the intention that the sample represents the whole" [Cox, 2015a]. Stasch et al. [2011a] proposes a Web Processing Service (WPS) that retrieves sensor data from a SOS service in order to aggregate it based on features-of-interest. The approach by Stasch et al. [2011b] is similar, but takes sensor data as input that is already published on the semantic web.

Ganesan et al. [2004] stresses that spatio-temporal irregularities are fundamental to sensor networks. Irregular sampling can have a potentially large

influence on the accuracy of the aggregated outcome. For example, averaging sensor data from a feature-of-interest that is being sampled densely in some parts and more sparsely in other parts could lead to inaccurate results. To counter this the values of the densely sampled area should have a lower weight than the values from the sparsely sampled area. The same holds true for temporal irregularities [Ganesan et al., 2004]. Also, Stasch et al. [2014] argue that in order for automatic aggregation to work there needs to be semantics on which kind of aggregation methods are appropriate for a specific kind of sensor data. Not all kinds of aggregation are meaningful (e.g. taking the sum of temperature values). This requires a formalisation of expert knowledge which they call semantic reference systems.

3 | METHODOLOGY

A number of studies related to this thesis have been reviewed in Chapter 2. This chapter discusses why the semantic web will be used for linking sensor metadata and which methods will be used to achieve this. The SWE standards, the om-lite and sam-lite ontologies, and RDF will be described. Also a brief overview of the data use for this thesis is presented.

3.1 SENSOR METADATA ON THE SEMANTIC WEB

Sem-SOS [Henson et al., 2009; Pschorr, 2013] as well as SEL [Janowicz et al., 2013] focus on combining the sensor web with the semantic web, but do not address the integration and aggregation of sensor data. Similarly, Atkinson et al. [2015] proposes to expose sensor data to the semantic web in order to find other kinds of related data about the same feature-of-interest. Data that can be collected for another area of research. However, Atkinson et al. [2015] do not mention the integration of complementary sensor data from heterogeneous sources either. Stasch et al. [2011b] and Stasch et al. [2011a] suggest interesting methods for aggregating sensor data based on features-of-interest. However, also these studies use sensor data from only a single source into account. Moreover, Corcho and Garcia-Castro [2010] and Ji et al. [2014] argue that methods for integration and fusion of sensor data on the semantic web is still an area for future research. Data fusion is "a data processing technique that associates, combines, aggregates, and integrates data from different sources" [Wang et al., 2015a, p. 2].

Jirka and Nüst [2010] and Jirka and Bröring [2009] present methods for including SOS services in an OGC catalogue service using SOR and SIR. Making sensor metadata available in a catalogue service will improve the discovery. However, discovery through the semantic web is likely to be more effective, since links can be created towards the sensor data from many different sources of related information. Another advantage is that links can be created by everybody that publishes linked data on the web, allowing sensor data to be used for implementations that were not identified beforehand by the publisher. Also, the semantic web will be easier to access, while the catalogue service can only be requested at a certain URL which has to be known to potential users.

Since data on the web has a distributed nature it can be questioned whether centralised catalogue services are feasible to create. It places a burden on the owner of the SOS to register with a catalogue service. Also, there could be multiple of these services on the web creating issue regarding the discovery of relevant catalogues. The semantic web could solve this issue by getting rid of the 'dataset-centric' approach and adding metadata directly to the web instead.
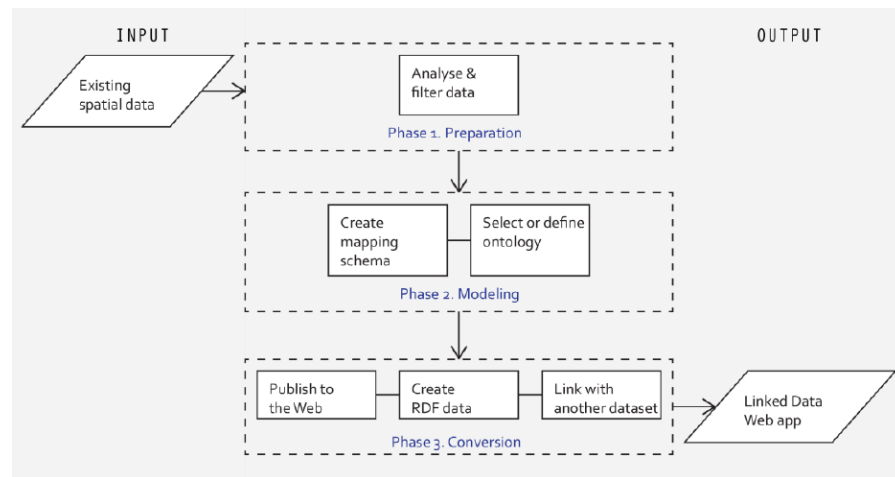
Figure 3.1: Workflow diagram for publishing linked open data from existing spatial data [Missier, 2015, p. 28]

3.2 CREATING LINKED DATA

For publishing sensor data on the semantic web a conversion of this data to RDF is required. For this the method by Missier [2015] will be used. She developed a workflow for publishing linked open data, using existing spatial data as input (Figure 3.1). This workflow consists of three phases: preparation, modelling and conversion. The next paragraphs will describe these three phases in more detail.

3.2.1 Preparation

The preparation phase deals with data acquisition, analysis and filtering. After an acquired dataset has been selected to convert to linked data it should be carefully examined. When creating linked data it is important to know the content and format of the input data. Missier [2015] explains that this understanding of the data is crucial for selecting the right ontologies and using the right software tools to process the data. Data filtering should be performed to select the parts of the dataset that need to be mapped to linked data. In this step the data quality could also be improved if necessary. The result of this phase should be a clean dataset, with semantics about the content, which has been filtered to only contain the parts of data that are required as linked data.

3.2.2 Modelling

The first step in modelling data to linked data is to select an ontology. An ontology is part of a data model, that contains semantics of how objects in the real world are mapped inside a dataset. To improve interoperability is preferred to (re)use an already existing ontology. However, if there is no suitable ontology available a new one should be created. To select an ontology all ontologies that describe the type of data inside the dataset should be listed. The one that fits the dataset and the final application best should be selected from these ontologies. Missier [2015] stresses that it is also important to reuse common predicates associated with an ontology. This makes

the data more understandable and it is easier to merge with other datasets using the same ontology and predicate.

3.2.3 Conversion

In the conversion of data to linked data Missier [2015] describes two approaches: the RDF storage approach and the 'on-the-fly' conversion. In the first approach data is converted to linked data and stored in a triple store. This triple store can be queried by a SPARQL endpoint. Alternatively, data can be stored in a spatial database and converted to RDF on the fly when it is requested. There are advantages and disadvantages to both methods. In general creating a triple store creates duplicate data, which should be very well maintained to prevent the duplicate data being out of sync. On the other hand creating a triple store is a one time operation after which the querying with SPARQL is quite efficient. The on-the-fly conversion does not create duplicate data, but has to convert its data to RDF for every SPARQL query.

3.3 SPATIAL QUERIES WITH SPARQL

The SPARQL query language can perform many different kinds of queries on RDF triples. However, in SPARQL no spatial queries have been implemented. For this purpose GeoSPARQL and stSPARQL have been created. The next paragraphs will describe these two extensions of SPARQL in more detail.

3.3.1 GeoSPARQL

GeoSPARQL "defines a vocabulary for representing geospatial data in RDF, and it defines an extension to the SPARQL query language for processing geospatial data" [Perry and Herring, 2012, p. xvi]. It allows for defining geometric data in RDF and performing spatial queries. The Dimensionally Extended Nine-Intersection Model (DE-9IM) [Strobl, 2008] has been implemented to find topological relations between two geometries. GeoSPARQL has been implemented in the 'Parliament' SPARQL endpoint [Battle and Kolas, 2012].

Listing 3.1: A GeoSPARQL query to find the names of features that contain a point geometry

```
PREFIX geo: <http://www.opengis.net/ont/geosparql#>
PREFIX geof: <http://www.opengis.net/def/function/geosparql/>
PREFIX foaf: <http://xmlns.com/foaf/0.1/>

SELECT
?name
WHERE {
?feature geo:hasGeometry ?geom .
?feature foaf:name ?name.
FILTER
    (geof:sfContains(?geom,"<http://www.opengis.net/def/crs/EPSG/0/4258>
    POINT(4.289244 52.027337)"^^geo:wktLiteral))
}
```

3.3.2 stSPARQL

The Strabon endpoint uses stRDF, which is "a constraint data model that extends RDF with the ability to represent spatial and temporal data" [Koubarakis and Kyzirakos, 2010, p. 425]. The stRDF model can be queried using stSPARQL, which syntax is similar to GeoSPARQL (listing 3.1 & 3.2). Both extensions of SPARQL use filter expressions to perform spatial operations on WKT or GML geometries. The definition of geometries and the syntax of the filter expression differ slightly.

Listing 3.2: A stSPARQL query to find the names of features that contain a point geometry

```
PREFIX strdf: <ttp://strdf.di.uoa.gr/ontology#>
PREFIX foaf: <http://xmlns.com/foaf/0.1/>

SELECT
?name
WHERE {
?feature strdf:hasGeometry ?geom .
?feature foaf:name ?name.
FILTER (?geom contains "POINT(4.289244
    52.027337);<http://www.opengis.net/def/crs/EPSG/0/4258>"^^strdf:WKT)
}
```

3.4 ONTOLOGIES

When publishing data on the semantic web, ontologies are required to specify what things are and how they relate to other things. The evaluation of observation metadata ontologies by Hu et al. [2014] is interesting, since it exposes what the relevant aspects are in the process of observation discovery. However, their proposed model focusses mainly on including remote sensing and imagery data in metadata models that were not originally created for this kind of data. The SSNO is an ontology that clearly describes the process between sensor, stimulus and observation. However, Cox [2015b] points out that an important aspect of describing a sensor network is missing in this ontology: the sampling. Also, the om-lite and sam-lite ontologies by Cox [2015b] are lightweight ontologies that can be complemented by already existing linked data ontologies. They do not rely on the (heavy) ISO specifications that date from before the semantic web, unlike the SSNO. The om-lite and sam-lite ontologies will therefore be used in this thesis.

The SOS has a number of metadata attributes such as the service provider's details (including contact information), its spatial and temporal extent (spatialFiler & temporalFilter) and the capabilities to query a subset of this extent. It receives data from a sensor which makes observations. An observation can be defined as "an action whose result is an estimate of the value of some property of the feature-of-interest, obtained using a specified procedure" [Cox, 2015a].The sensor is placed at a sampling point. The sampling point is part of a sampling feature which intents to resemble the feature-of-interest. In the case of air quality the feature-of-interest is the bubble of air surrounding the sensor, therefore the sampling point equals the feature-of-interest [INSPIRE, 2014]. The design is that an observation of the sampling feature describes the feature-of-interest through measuring one of its prop-

Figure 3.2: Artifacts of the WPS service model [Open Geospatial Consortium, 2015, p. 15]

erties. The measurement procedure is described by a short string of text, input and output parameters and the units of measurement of the ouput. The relation between feature-of-interest and administrative units is added to improve the discovery of sensor data on the semantic web.

To publish data on the semantic web ontologies are required to specify the different classes and their relations. An ontology for static geographic data has to be connected to an ontology for sensor metadata. From the UML diagram in Figure **??** the classes Observation, Process, ObservedProperty and FeatureOfInterest can be mapped to classes belonging to OWL for observations [Cox, 2015c]. SamplingFeature and Sampling point can be mapped to classes from OWL for sampling features [Cox, 2015d]. GeoSPARQL can be used for the administrativeUnit class [Perry and Herring, 2012] and the PROV ontology for the sensor and sensor observation service classes [W3C Semantic Sensor Network Incubator Group, 2011].

3.5 WEB PROCESSING SERVICE

The OGC Web Processing Service is a standard interface for making simple or complex computational processing services accessible as web services. Originally, it has been created with spatial processes in mind, but it can also be used to insert non-spatial processing into a web services environment [Open Geospatial Consortium, 2015, p. 8]. Via the WPS jobs can be controlled and monitored, which run certain processes (Figure 3.2). Similar to WFS, WPS and SOS, the WPS has requests for retrieving metadata: `GetCapabilities` and `DescribeProcess`. On top of that there are is the `Execute` request to execute a process, and a number of other requests for various purposes: `GetStatus`, `GetResult` and `Dismiss`. All requests will be briefly described in this paragraph.

3.5.1 Get Capabilities

All OGC web services give an overview of what they have to offer using the so-called `GetCapabilities` request. The request can be made by taking the HTTP address of the WPS and adding: `service=wps&request=getcapabilities`.

Figure 3.3: UML model of WPS capabilities document [Open Geospatial Consortium, 2015, p. 70]

This returns a document with information about the service metadata, the basic process offerings, and available processes. Figure shows the model of this capabilities document. The service identification section of the document gives a description of the service, including the versions that it supports and potential fees or access constraints. The service provider section gives information about the organisation or people who maintain the WPS and also includes their contact information. The operations metadata lists the different requests that are implemented in this particular WPS and the HTTP addresses to which the GET and POST requests can be send to. The content section of the capabilities document provides an overview of the available process offerings. For every process an identifier, title and description are listed. Optionally an extensions section can be added that describes additional service capabilities. At the end of the document the language section lists the languages that are supported and the default language that is being used. An example of an capabilities document can be found in Appendix B.1

3.5.2 Describe Process

A more detailed description of a process listed in the capabilities document of the WPS can be retrieved using the `DescribeProcess` request. This request requires the identifier of the process to be passed as a parameter. Optionally, a specific language can be requested for the response document (from the list of available languages in the capabilities document). The process description includes information about input parameters and output data, such as their identifiers, data type, mime types or default values. A describe sensor request can be made by putting the base address of the WPS and adding: `service=wps&version=1.0.0&request=describeprocess&identifier=an_identifier` where the version parameter should be in accordance with the supported version(s) from the capabilities document and the identifier should be taken from the process offerings

section of the capabilities document. An example of an describe sensor response document can be found in Appendix B.2

3.5.3 Execute

A WPS process can be started using the Execute request. To make this request the HTTP address of the WPS is extended with: `service=wps&version=1.0.0&request=execute&identifier=GetSensors` where the version parameter should be in accordance with the supported version(s) from the capabilities document and the identifier should be taken from the process offerings section of the capabilities document. Additionally, the desired execution mode can be added to the request (synchronous, asynchronous or auto) as well as the desired output format (response document or raw data). By default the execution mode is set to 'document'.

Depending on the individual requirements of the process, input parameters can be added using `&DataInputs=[parameterName1=value1; parameterName2=value2]`. The parameter names are defined in the describe process response document, as well as the allowed values for each parameter. There are two kinds of input parameters that can be put in a Execute request: literal and complex data inputs. Literal data inputs can be a string of characters in consisting of different data types (e.g. Double, Integer, String), a given value range, and associated units (e.g. meters, degrees Celsius) [Open Geospatial Consortium, 2015, p. 36].

Complex data inputs are made for inputting complex vector-, raster- or other kind of data. This data can be inserted directly in the request or indirectly by referencing to a file. The process will then first retrieve this file from a remote server before running. A complex data input defines the allowed mime types that the process accepts.

When the process finishes an execute response document is retrieved. This document has a process section with the identifier, title and abstract of the finished process. It also contains a status section with the time the process ended and whether it finished successfully. If any output data has been produced an 'ProcessOutputs' section is created that contains the identifiers of the outputs and the corresponding data. An example of an execute response document can be found in Appendix B.3

3.5.4 Other requests

WPS processes can run synchronously or asynchronously. With a synchronous execution the connection with the client stays open until the process has finished. However, for processes that take longer to execute the asynchronous mode is better suited. The process will continue running after the connection has been closed. With a GetStatus request the client can check whether the process is still running. This request is structured the same as the execute request, but with the mode set to 'status'. Once it has finished the GetResult request allows the client to retrieve the output data. The Dismiss request can be made to communicate to the server that the client is no longer interested in the results of a job. This job will then be cancelled and its output deleted. A job identifier is a required parameter for all three of these requests.

3.6 DATA

A number of datasets have been used in this thesis. The following paragraphs will go over the different datasets, describing their sources and contents.

3.6.1 Vector data

Topography

The datasets of Dutch provinces (provincies, Figure A.2) and municipalities (gemeenten, Figure A.1) have been downloaded from `https://www.pdok.nl/nl/producten/pdok-downloads/basis-registratie-kadaster/bestuurlijke-grenzen-actueel`. For the Netherlands there are 12 features in the provinces and 393 in the municipalities dataset.

It has been challenging to obtain data of administrative boundaries of Belgium (even from the INSPIRE data portal). Therefore, all data for Belgium was retrieved from `http://www.gadm.org/`. There are also 12 features in the provinces (including the capitol region of Brussels) and there are 589 features in the municipalities dataset.

The country datasets have also been downloaded from `http://www.gadm.org/` (Figure A.3). The administrative unit data contains the names of the administrative units and their (polygon) geometry.

Land cover

Data on land cover will be used to complement the data of administrative units. A section of the 2012 dataset from the Coordination of Information on the Environment (CORINE) programme will has been selected for this (Figure A.4). The entire CORINE dataset was retrieved from `http://land.copernicus.eu/pan-european/corine-land-cover/clc-2012`. The features overlapping the Netherlands and Belgium have been retrieved from this dataset using the open source QGIS software and stored in a separate database in Postgres.

The database contains polygon geometries (Figure A.5) with a unique identifier and a code that refers to the type of landcover. These codes can be looked up in the accompanied spreadsheet file containing the legend table of CORINE 2012.

3.6.2 Raster data

Data is often used in a raster representation for computations in a Geographical Information System (GIS). For natural phenomenon a raster representation is especially well suited. The European Environment Agency (EEA) reference grid is a standard grid which covers Europe. It is available with a resolution of 100km², 10km² and 1km². In this thesis the EEA grid cells with a resolution of 100km² (Figure A.6) and 10km² (Figure A.7) have been used that overlap the Netherlands and Belgium. 15 grid cells of 100km² and 843 grid cells of 10km² have been selected from the original dataset.
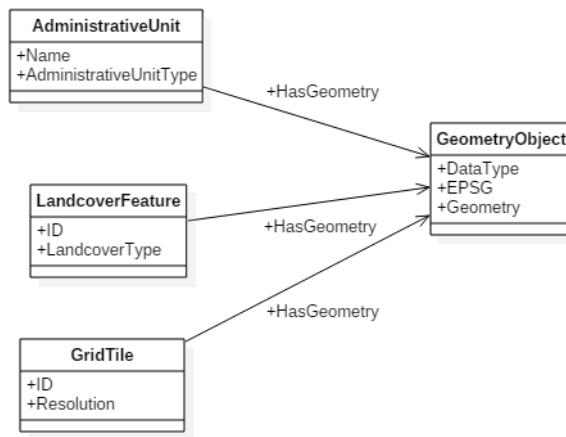
Figure 3.4: Model of vector and raster features

3.6.3 Sensor data

Air quality sensor data will be used from the RIVM (http://inspire.rivm.nl/sos/) and from the Belgian interregional environment agency (IRCEL-CELINE) (http://sos.irceline.be/). Both of these organisations have a SOS where data can be retrieved according to the SWE standards. The one of the RIVM has been online since the 21st of August, 2015. IRCEL-CELINE already made the SOS available on the first of January, 2011. Figure A.8 and Figure A.9 show the sensor networks of both organisations. They provide different kinds of sensor data, such as particulate matter (PM_{10}), nitrogen dioxide (NO^2) and ozone (O^3). Figure A.10 shows one of the sensor locations in the city center of Amsterdam.

3.7 PREPARING LINKED DATA

Linked data has been prepared that is used to retrieve and process sensor data on the semantic web (Figure 3.4). This is done for vector data sets of administrative units and land cover features, and for raster data sets of EEA grids with a resolution of 10km² and 100km².

Three types of administrative units have been converted to linked data: countries, provinces and municipalities. Every administrative unit has a name, 'type' and (multi)polygon geometry assigned to it (Figure 3.4). The administrative unit type is defined by DBPedia URIs of country, province and municipality.

The CORINE 2012 land cover dataset contains features with an identifier, a land cover type and a (multi)polygon geometry (Figure 3.4). The identifier has the form of: 'EU-' plus a unique seven digit number. The land cover type is defined by a three digit number, which can looked up in the provided spreadsheet containing the legend.

The EEA reference grid with resolutions of 10km² and 100km². Every feature is defined by an identifier, a resolution and a point geometry of the origin (Figure 3.4). The identifier is a code given to a feature by the EEA and has the form of: resolution + 'E' + x coordinate + 'N' + y coordinate.

4 | DESIGN

This thesis aims to design a method that uses the semantic web to improve sensor data discovery as well as the integration and aggregation of sensor data from multiple sources. The outline of this method will be described in this chapter using the existing methods from Chapter 3.

4.1 CREATING LINKED DATA FROM SENSOR META– DATA

The process of automatically creating linked data from sensor metadata is shown in Figure 4.1. A WPS contain processes for retrieving metadata, converting it to linked data and outputting it to a triple store. Data from a SOS is retrieved by a WPS process. This process converts it to linked data. The output is an RDF document containing the metadata as triples. These documents are posted to a SPARQL endpoint, where they can be queried.

The workflow of this WPS process is shown in Figure 4.2. It is an adaption of the workflow by Missier [2015], which was originally intended for creating linked data about vector parcel data. The input of the process should be the HTTP address of a SOS. Since both the requests and the data model are standardized in a SOS the process should be able to automatically perform the tasks for creating linked data. The first step is to make requests to the SOS to retrieve it's metadata. This results in a number of XML documents that need to be filtered. The second step is to map the data inside these XML documents to linked data ontologies. In the final step RDF documents are created of the mapped metadata and published on the web.
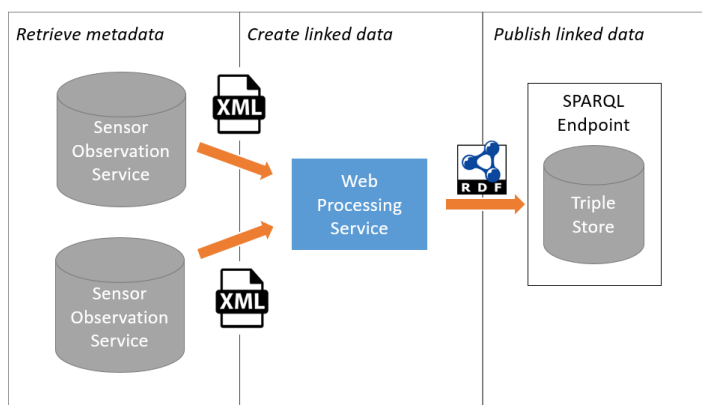
Figure 4.1: General overview of creating linked data of metadata from Sensor Observation Services

Figure 4.2: Workflow diagram of Web Processing Service (adapted from Missier [2015])

4.1.1 Retrieving metadata from the Sensor Observation Service

The first step of creating an online knowledge base with sensor metadata is to retrieve the metadata from the different Sensor Observation Services. This data has to be understood in order to map it to an ontology and it should be filtered to only contain the required parts of data. The next paragraphs will describe the way sensor metadata is modelled in a SOS, with which requests it can be retrieved and how it should be filtered.

Sensor metadata model

A sensor observation service describes a number of its properties that are required to know in order for clients to request data from it. It identifies the organisation that maintains it, with at least the organisation's name and its contact information. Optionally, the organisations website, keywords and an abstract about the SOS can be supplied. The SOS also describes its identifier and HTTP addresses (the address for sending requests can differ for POST or GET requests). It also lists the SOS versions and response formats it supports. The access constraints and fees are also mentioned. In most cases the use is free of charge and without access constraints. However, it is possible for an organisation to restrict the use of the SOS in these ways.

In the SWE standards a sensor is modelled using two entities: a procedure and a Feature of Interest (FOI). The procedure is the method of sensing and the FOI is the feature of which the sensor is sensing a certain property. Therefore, the observable property ties together the procedure and feature of interest. It should be noted that the geometry of a FOI is not necessarily always a point geometry. It can also be generalized into larger features (e.g. multiple sensors observing different parts of one lake).

In version 2.0 of the Sensor Observation Service Interface Standard [Bröring et al., 2012] an offering is defined as a grouping of FOIs, which have a common procedure. The constraint of sharing the same procedure has been added in version 2 of this document to solve the ambiguity of offerings in SOS 1.0. The purpose of offerings is to allow users to query the observation data more efficiently. FOIs that are often queried together are therefore grouped into the same offering for efficient retrieval.

Figure 4.3: Sensor metadata derived from a SOS

Metadata Requests

To retrieve metadata from a SOS a `GetObservation` request is made first. This is a request with a very generic structure. The GET request is created by adding `service=SOS&request=GetCapabilities` to the HTTP address of the SOS. For example, the RIVM has its SOS at the address: `http://inspire.rivm.nl/sos/eaq/service?`. Therefore, the capabilities document can be retrieved using the following URL: `http://inspire.rivm.nl/sos/eaq/service?service=SOS&request=GetCapabilities`.

This requests returns the capabilities document of the SOS (see Paragraph 2.1.3). This lists the identifier of each FOI, each procedure and each observed property. It also has a section where the offerings that it contains are being described. This description of an offering includes a unique identifier, a procedure and the corresponding observed property. Additionally, descriptions can be added such as a bounding box, temporal range, FOI type and response format.

Unfortunately, the capabilities document is not able to provide information about which procedure is being applied for which feature of interest. This is crucial information for knowing which deployed sensors can be queried using a particular SOS. Also, the features' geometries cannot be retrieved from the capabilities document. Based on that document it is not yet clear which sensor locations are being used and what is measured at a specific location. Therefore, a `GetFeatureOfInterest` request can be made to retrieve the location of each FOI. Such a request can be made by adding `service=SOS&version=2.0.0&request=GetFeatureOfInterest`. The version KVP should correspond to the version declared in the capabilities document. A pointer to a specific FOI is optional and usually all FOIs are returned by default. Using the example of the SOS by the RIVM the `GetFeatureOfInterest` request looks like

this: `http://inspire.rivm.nl/sos/eaq/service?service=SOS&version=2.0.0&request=GetFeatureOfInterest`.

On the one hand, a `GetFeatureOfInterest` document does not necessarily provide information about the procedures that are related to a certain feature of interest. On the other hand, a `DescribeSensor` request does not always relate the process to a FOI either. However, a `GetObservation` requests return observation data grouped per feature of interest. Therefore, small amounts of data can be retrieved from each offering using `GetObservation` requests to link the FOIs to procedures and observed properties. When possible a temporal filter should be used to limit the data traffic. Using this method every procedure and offering can be related to a set of FOIs with their corresponding geometry. This represents the collection of sensor devices of which data can be retrieved by sending requests to the SOS.

Filter Metadata

The documents that are returned by the SOS contain a lot of information. In some cases the returned information can be limited by adding filters to the requests. However, not all SOS have supported all filters and not all unnecessary data can be filtered out. Therefore, the XML documents that are returned should often be filtered on the client side. In a XML document every element should be defined using a namespace. Often these prefixes are defined in the xmlns tag at the top of the document to refer to these namespaces. These namespaces and corresponding tags can be used to filter the response documents for the content that is required. It should be noted that there are multiple namespaces that could be used to define the same concept. However, the potential namespaces that can be encountered are restricted by the schema describing the content of a response document. This schema is usually referenced to in the start tag of the response document.

4.1.2 Modelling with the om–lite and sam–lite ontologies

After the metadata has been retrieved from the SOS and filtered (Figure 4.3) it has to be mapped to linked data ontologies. For this the om-lite and sam-lite ontologies are being used in combination with the PROV and GeoSPARQL ontologies. Figure 4.4 shows the which ontologies that can be used to describe the classes of Figure 4.3.

A SOS is modelled as an agent with a specific name, that acts on behalf of a certain organisation. The organisation, access constraints, fees, versions and response formats are properties of the SOS. Every sensor is described by a procedure and a certain feature of interest. The sensor class was not present in the model of Figure 4.3, because the SOS does not define sensors. However, the sensor class has been added to the semantic model to make the relation between procedure and sampling point explicit.

In Figure 4.4 the collection of sampling features is added. These are collections of all FOIs from different Sensor Observation Services of which the same property is being observed. The sampling collections are currently modelled to only contain sampling points. This is because the FOI of an air quality sensor is equal to the bubble of air directly around the sampling point. Other types of sampling features can be added when the application requires this. The offering class is modelled as a specialization of the collection of sampling features. It contains a subset of the sampling points that are part of the same offering at a particular SOS.

Figure 4.4: Sensor metadata as modelled in RDF (om-lite classes in yellow and sam-
lite classes in purple)

Table 4.1: Types of PURLs [Shafer et al., 2016]

PURL Type	Meaning	HTTP Shorthand
301	Moved permanently to a target URL	Moved Permanently
302	Simple redirection to a target URL	Found
303	See other URLs (use for Semantic Web resources)	See Other
307	Temporary redirect to a target URL	Temporary Redirect
404	Temporarily gone	Not Found
410	Permanently gone	Gone

The last class that has been added to the model as shown in Figure 4.4 with respect to the model from Figure 4.3 is the DBPedia observed property class. Every observed property that is defined in a SOS relates to a certain observed property as defined by DBPedia. Since SOS requests require their own identifiers as input the observed property class exists twice in the model: one as defined by the SOS and one as defined by DBPedia. For the same reason all sampling points, processes and offerings have a 'name' attribute in addition to their URI. These store the original (non-semantic) identifier that they were given by the SOS.

4.1.3 Output linked sensor metadata

Once the data has been retrieved and mapped to their corresponding classes in the ontologies RDF triples can be created to link the data together. These triples should be stored in files and posted to a SPARQL endpoint. The following two paragraphs will describe these steps in more detail.

Create RDF

For every mapped part of metadata from the SOS one or more triples are created. These triples consist of at least of URIs. However, preferably they are URLs that can be resolved to an RDF document that contains semantic information about what it represents. To do so every part of metadata is automatically assigned a Persistent Uniform Resource Locator (PURL). For this kind of URL to resolve a PURL server should be set up. This server performs one of the six tasks shown in Table 4.1 whenever a PURL is being retrieved. The structure of a PURL consists of the HTTP address of the PURL server, with a unique identifier attached to it. For example, `http://www.examplePURLServer.com/unique_identifier`.

Once the metadata is assigned PURLs the triples can be created. Since the metadata of the sensors is being returned by the SOS in a structured way links can be created between URLs of FOIs, observable properties, procedures and the other classes shown in Figure 4.4. After these triples have been created the linked data should be serialized to an RDF document. For this a specific notation has to be selected.

Publish RDF on the web

To store the linked data created under the previous paragraph a SPARQL endpoint has to be set up. This endpoint is connected to a triple store. This means that the RDF documents containing linked data can be send to it using POST requests. After it has been stored in the triple store users can query the triples using the SPARQL query language.

Once the linked data is inside the triple store the PURLs can be redirected to SPARQL `Describe` queries. A describe query takes a URI as input and

returns all the triples that contain this URI as either an subject, predicate or object. This means that all information that the endpoint has about a particular URI is being returned. If the PURL server receives a request for a particular PURL it makes a `Describe` query to the endpoint and returns to the client all the linked data that it retrieved about this PURL.

4.2 USING LOGICAL QUERIES TO RETRIEVE SENSOR DATA

The second web process looks on the semantic web for sensors that observe a certain property in a specific area. It collects the data for these sensors at their corresponding Sensor Observation Service. When multiple data sources are found the data is integrated into a single dataset. The sensor data is temporally aggregated before it is returned to the user. Optionally, spatial aggregation can be performed as well.

4.2.1 Discovering sensors

For discovering sensors there are a number of input parameters for the second process: An observed property, a set of names of spatial features, a temporal range and granularity. Additionally, a type of spatial aggregation can be added. The input data is inserted in a number of SPARQL queries. First the geometry is retrieved for all features inside the input list of features. This is done using the name or identifier attribute of spatial features at the endpoint.

Then a second SPARQL query selects the sensors that are part of a collection that observe a certain property of the FOIs. This query also contains a spatial filter with the retrieved geometries. This makes sure only sensors in the requested areas are returned. For all discovered sensors the corresponding SOS URL, observed property, offering and process identifiers are retrieved. These original SOS identifiers are required for creating `GetObservation` requests.

4.2.2 Retrieving sensor data

The SPARQL queries return a table with all necessary information per sensor on each row. The columns of this table are: sensor URI, FOI geometry, FOI identifier, procedure identifier, observed property identifier, offering identifier, and the last column contains the URL of the SOS. The data that is returned by the SPARQL queries described in Paragraph 4.2.1 correspond to the `GetObservation` input parameters as described in Paragraph 2.1.3. This request is send out using all the values per row in the returned table of sensor metadata and is structured by taking the URL of the SOS and extending it with:

```
service=SOS&version=2.0.0&request=GetObservation&procedure=
the_procedure&offering=the_offering&observedproperty=the_
observed_property&responseformat=http://www.opengis.net/
om/2.0&featureOfInterest=the_feature_of_interest.
```

If the SOS supports temporal filters these should be added to only retrieve observation data from the temporal range that the user is interested in. To

include a temporal filter the following parameters can be added to the above GET request:

```
&temporalFilter=om:resultTime,start_time/end_time
```

4.2.3 Integrating and aggregation sensor data

The requests described in Paragraph 4.2.2 result in a XML document with observation data for every sensor. In each document the observations have to be retrieved, with their corresponding result time, phenomenon time and UOM. All of these observation have to be stored in a uniform way in another file or file-like object. If the temporal filter was not supported by the SOS the observation should only be stored if it is inside the requested temporal range. Also the FOI geometry should be stored with the observation for data visualisation or spatial aggregation later on. This process repeats itself until the observations have been retrieved from all XML documents and they are all stored in a single dataset.

The next step is to temporally aggregate the observation data. For every sensor location the result time is compared to the temporal range and the value for temporal granularity. It should then be appended to a list that corresponds to a certain subset of the temporal range. After this temporal sorting has been performed all observation values in a list should be aggregated according to a user defined method.

If a spatial aggregation method is part of the user's logical query then this is performed after the temporal aggregation. For each of the spatial features for which observation data has been retrieved a single list of aggregated data is produced. This list contains the observation value for each time interval (with a length equal to the temporal granularity) inside the temporal range. The spatial aggregation is performed in a similar fashion as the temporal aggregation. First all observation data is order in a list per time interval per spatial feature. After all observations have been ordered the required aggregation method is performed to produce a single value for each of the lists.

Examples of aggregation methods that can be used for spatial or temporal aggregation are the average, median, minimum, maximum or sum. Which method is being used is dependent on the information need of a user. Therefore, these methods can be extended with more complex aggregation techniques as described by [Ganesan et al., 2004]. Also, as Stasch et al. [2014] argue a 'check' could be implemented to make sure users do not (accidentally) request a meaningless type of aggregation, such as the sum of temperature values over a certain area.

5

PROTOTYPE IMPLEMENTATION

The prototype implementation serves as a proof of concept. It takes the designed method from Chapter 4 and applies it using existing software tools and python libraries and modules. First, the implementation of the process for automatically creating linked sensor metadata will be described. After that the implementation of the process for retrieving sensor data using logical queries is described. The third part of this chapter describes how the two processes have been made available online. The final part shows the web application that has been created as an interface for the second web process.

5.1 CREATING LINKED DATA FROM SENSOR META–DATA

The first step of the method described in Chapter 4 is to automatically harvest sensor metadata from a SOS using various requests. This data should then be mapped to ontologies and serialized to an RDF document. The next paragraphs will describe how these steps have been implemented using the Python programming language.

5.1.1 Making sensor metadata requests

A Python class object is created for the SOS based on Figure 4.3. This class contains the different variables and has built in functions to automatically retrieve the metadata. When a SOSclass instance is created the URL of the SOS has to be entered as input value. The initialisation of the SOSclass instance creates empty variables are created for storing information about the SOS' organisation, supported SOS versions, response formats, FOIs, offerings, observable properties and procedures. The FOIs are stored in a Python dictionary to allow the geometry to be stored, as well as information about the Coordinate Reference System (CRS), and the procedures and observed properties it is related to. Similarly, the procedures are stored in a dictionary to be able to easily access which observable property, which FOIs and which offerings are related to a procedure.

After the initialisation is finished the method `SOSclass.request()` is automatically triggered. This function starts by sending a `GetCapabilities` request to the SOS. For making HTTP GET and POST requests the Requests library for Python is used (see http://docs.python-requests.org). Listing 5.1 shows how to create a `GetCapabilities` request with the HTTP GET function from this library.

Based on the content that is being retrieved from the capabilities request further requests are being made. First, the geometries of the FOIs are collected using `GetFeatureOfInterest` requests. Similar to the `GetCapabilities` there are no specific parameters required except for:

`service=SOS&version=2.0.0&request=GetFeatureOfInterest`. However, one of the Sensor Observation Services used in this thesis had implemented their `GetFeatureOfInterest` request to always require a feature id parameter. Therefore, the following exception had to be built in case this error is returned: `service=SOS&version=2.0.0&request=GetFeatureOfInterest& featureOfInterest=allFeatures`.

After the geometries of the FOIs have been retrieved they should be linked to procedures and observable properties. As described in Paragraph 4.1.1 it is not mandatory to define these relations in the `GetCapabilities`, `GetFeaturesOfInterest` or `DescribeSensor` response documents. Therefore, small amounts of observation data are requested from the sensors using `GetObservation` requests.

Listing 5.1: Creating a HTTP Get request using Python's Request library

```
import requests


sosURL = "http://example.com/SOS?"


# Create the GetCapabilities request string
GetCapabilities =
    "{0}service=SOS&request=GetCapabilities".format(sosURL)


# Send the request
r = requests.get(GetCapabilities)


# Print the response document
print r.content
```

5.1.2 Map metadata to ontologies

After each request is send an XML document is returned. To retrieve data from these XML documents the LXML library for Python is used (see `http:// lxml.de`). With this library the XML document can be loaded into an Python object, which allows for easy XML processing, such as looping through the elements and searching the whole document for elements with specific tags. Listing 5.2 shows a snippet of code that takes the response document retrieved from Listing 5.1 and that uses the LXML library to find all offerings presented in this document. All offerings are returned as a list and stored inside the variable 'SOSclass.offerings'. With this principle all metadata from the SOS is retrieved from the XML response documents and stored for further processing.

Listing 5.2: Creating an Etree object from an XML response document using Python's LXML library

```
import lxml


# Store the retrieved document as an Etree object
tree = etree.fromstring(r.content)


# Retrieve the namespaces from the XML document
nsm = tree.nsmap


# Find all subsets of the XML document that are inside a
# 'sos:ObservationOffering' element
```

```
SOSclass.offerings = tree.findall(".//sos:ObservationOffering", nsm)
```

Once all the relevant metadata has been retrieved from the SOS and stored inside the class object it should be mapped to linked data ontologies. For the Python package RDFlib is used (see https://rdflib.readthedocs.org/). RDFlib defines an RDF graph to which triples can be added. Listing 5.4 shows a snippet of code that defines an RDF graph and adds all procedures of a SOS to it with the type 'http://def.seegrid.csiro.au/ontology/om/om-lite#process'. A semantic URL is defined for each procedure. It is created using a combination of the name of the organisation and a unique number for each procedure. The domain of the URL points to the PURL server. The same principle of creating triples from sensor metadata is applied to all classes and relations described in Figure 4.4.

Listing 5.3: Creating an RDF graph object with the Python package RDFlib

```
import rdflib

# The domain of the PURL server
PURLZ = "http://localhost:8099/masterThesis_tudelft"

# Create the OM-lite namespace
om_lite =
    rdflib.Namespace('http://def.seegrid.csiro.au/ontology/om/om-lite#')

# Initialize a graph object
g = Graph()

for i, procedure in enumerate(SOSclass.procedures):
  # Define a URIs for the procedures
  procedureURI = URIRef("{0}/{1}_PROC_{2}".format(PURLZ,
      SOSclass.organisation, i))

  # Add all procedures to the graph and define them
  # as om-lite processes
  g.add( ( procedureURI, RDF.type, om_lite.process) )
```

5.1.3   Publish linked data

For publishing linked data the Strabon endpoint (Figure 5.1) is used in combination with an Apache Tomcat server. The Strabon endpoint is a semantic spatiotemporal RDF store, originally developed for the European SemsorGrid4Env project (see http://strabon.di.uoa.gr/). It uses a Postgres database with the Postgis extension to store RDF triples and it provides spatial queries using stSPARQL (Paragraph 3.3.2).

The first step in publishing the graph that is created using RDFlib is to store it in an RDF document. This process is called the serialization of an RDFlib graph. To perform this process function 'serialize' of the RDFlib package is used. This function is a method of the graph object and requires two input parameters: the name of the output file and the notation of the triples. In the prototype implementation the Turtle notation is being used.

After the RDF documents have been created they should be posted to the SPARQL endpoint. To post a document to Strabon the client should first login to the endpoint. To do this a Session object is created using the Requests library. This session object posts it login credentials to the endpoint, after

Figure 5.1: User interface of the Strabon endpoint

which it can be used for posting RDF data. Listing 5.4 shows how to log in to the endpoint with a session object. After that the graph is serialized to the document 'sensors.ttl', which can then be posted to the Strabon endpoint.

Listing 5.4: Serializing the RDFlib graph object and posting it to the Strabon endpoint

```
import os
import requests
import rdflib

# Define the login parameters in a dictionary
login = {"dbname":"endpoint", "username":"your_username",
    "password":"your_password", "port":"5432",
    "hostname":"localhost", "dbengine":'postgis'}

# Create a new session
session = requests.Session()

# Log in with the session to the Strabon endpoint using the login
    parameters
r = session.post("http://example.com/strabon-endpoint/DBConnect",
    data=login)

# Serialize the graph to the file sensors.xml
g.serialize("sensors.ttl",format="turtle")

# Post the sensors.xml file to the endpoint
r = session.post("http://example.com/strabon-endpoint/Store",
    data={"view":"HTML", "format":"Turtle",
    "url":"file:///{0}/sensors.ttl".format(os.getcwd()),
    "fromurl":"Store from URI" })
```

For creating Persistent Uniform Resource Locators the Purlz software has been used (see http://www.purlz.org/). All URIs that are created get a PURL assigned to it. The PURL resolves the URI to a DESCRIBE query at the endpoint. This query is structured as a get request: `http://localhost /strabon-endpoint-3.3.2-SNAPSHOT/Describe?submit=describe&view =HTML&handle=download&format=turtle&query=DESCRIBE<an_URI>`. The request has '/Describe?submit=describe' to call the script that deals with describe queries and to tell it that the request is also submitting this type of query. The parameters 'view=HTML&handle=download' indicate that the endpoint's website is requested, but the returned data should be a download file instead of an HTML page. The parameter '&format=turtle' sets the RDF notation of the download file to Turtle and '&query=DESCRIBE <an_URI>' is the SPARQL query that contains the URI between brackets.

Listing 5.5: Example of a PURL batch file (containing one PURL)

```
<purls>
  <purl id="/masterThesis_tudelft/responseFormat_0" type="302">
    <maintainers>
    <uid>admin</uid>
    </maintainers>
    <target url="http://example.com/strabon-endpoint/Describe?
      view=HTML&amp;handle=download&amp;
      format=turtle&amp;submit=describe&amp;query=DESCRIBE
      &lt;http://localhost:8099/masterThesis_tudelft/RIVM_PROC_0&gt;"/>
  </purl>
</purls>
```

Every URI that is assigned to a part of sensor metadata (like the 'proce-dureURI' in Listing 5.4) is written to an XML file with the parameters: ID, PURL type, and target address. Optionally, information about the person or organisation maintaining the PURL can be added. The ID is the original URI that is being resolved to the target address. The PURL type is set to 302, which means that it redirects the client to the target address. Alternative types can be found in Table 4.1. After all URIs have been added to the so called XML 'batch' file [PURL, 2016], the file can be posted to the Purlz server. Posting XML documents to the PURL server is similar to posting RDF documents to the Strabon endpoint as shown in Listing 5.4. An example of a 'batch' file can be seen in Listing 5.5. These XML files are created using the Python package LXML.

5.2 USING LOGICAL QUERIES TO RETRIEVE SEN-SOR DATA

The first process described how the prototype implementation harvests sensor metadata from a SOS and publishes it as linked data in a SPARQL endpoint. The second process uses this data and its semantics to automatically retrieve data based on a user's logical sensor data query. This process is created around 'request' class which has been made to store all logical input parameters and automatically convert them to SOS requests and return the data in an integrated and aggregated way. The following paragraphs describe how the different parts of this process have been implemented.

5.2.1 Input parameters

The prototype implementation takes a number of input parameters. First of all, the observable property, which will be related to a DBPedia definition. The second parameter is the category of input features. This can be set to administrative units (country, province or municipality), land cover or raster. The third parameter is a list of input feature. This is a list of names or identifiers that correspond to the category of input features. For example, 'Delft municipality' can be used as input since it includes the feature name and it's category. The next parameter is the temporal range. This has to be a list of two ISO datetime strings representing the start and end time. The fifth parameter is the temporal granularity, represented by an integer and a datetime unit like hour(s), day(s) or week(s). For example a temporal range of '1 hour' will aggregate the data temporally to a single value per hour. The sixth input parameter is the temporal aggregation method. This is the method for aggregating data between start and end time using the provided temporal granularity. The last input parameter is the method of spatial aggregation. This method will be applied to aggregate the data based on the input features.

5.2.2 Discovering sensors

The input category is a starting point for the process to find the geometries of the input features. It creates a SPARQL query that retrieves the geometries of the features by selecting all features of the input category and filtering them by name. Listing 5.6 shows that the filter expression is defined first. This is then added to the a predefined template of a SPARQL query together with the feature category. For the currently implemented categories (landcover, raster, municipalities and provinces) the DBPedia URL is simple found by adding it to the standard method for defining URIs: http://dbpedia.org/resource/.

Listing 5.6: Example SPARQL query retreiving the geometries of input features

```
import requests

# Define the location of the endpoint
myEndpoint = 'example.com/strabon-endpoint/Query'

for i, feature in enumerate(self.featureNames):

  # Create a list with a filter expression for every name
  featureNamesList = '?name = "{0}"'.format(feature)

  # Join them together and place them inside the SPARQL filter
      function
  featureFilter = "FILTER( {0} )".format( " ||
      ".join(featureNamesList) )


# Define the SPARQL query
query = """
  SELECT
    ?feature ?geom ?name
  WHERE {{
```

```
    ?feature <http://www.w3.org/1999/02/22-rdf-syntax-ns#type>
        <http://dbpedia.org/resource/{0}> .
    ?feature <http://strdf.di.uoa.gr/ontology#hasGeometry> ?geom .
    ?feature <http://xmlns.com/foaf/0.1/name> ?name .
  {1}
  }}""".format(self.featureCategory.title(), featureFilter)

# POST the SPARQL query to the endpoint
r = requests.post(myEndpoint, data={'view':'HTML', 'query': query,
    'format':'SPARQL/XML', 'handle':'download', 'submit':'Query' })

# Print the results
print r.content
```

A SPARQL query is then made to find a sensor collection that has a re-quested observed property. From this collection sensors are selected that overlap with the previously found geometries. Unfortunately SPARQL queries are not allowed to exceed a certain number of characters. This creates prob-lems when querying larger vector geometries (provinces and countries). For these queries two alternatives have been implemented: using the EEA refer-ence grid as a spatial index for vector geometries and using bounding box queries at the SOS.

For the first alternative the EEA raster cells are retrieved instead of the vec-tor data. Only the cells are requested that overlap with the vector geometry. For these cells all sensor locations are requested. However, the result of this is that too many sensor locations are retrieved, also ones that are outside the original vector feature. Therefore the WPS performs the spatial filter and removes all locations that are outside of the requested feature.

The second alternative checks which Sensor Observation Services have sensor locations within the bounding box of the vector geometry. For all of them have sensors in this area that observe the required observed prop-erty `GetObservation` requests are made. In these requests a spatial filter is added.

5.2.3 Retrieving sensor data

After all sensor locations have been retrieved from the SPARQL endpoint `GetObservation` requests are made. These require the identifiers that were given to the observed properties and procedures by their SOS, instead of the semantic URLs that were assigned to them. The requests are structured as explained in Paragraph 2.1.3. When possible, the request is extended with a temporal filter to only retrieve data inside the required temporal range. The output format is set to XML using the http://www.opengis.net/om/ schema.

The received XML document loaded into Python using the LXML package and all 'sos:observationData' elements are retrieved from it. These elements contain O&M observations. Some Sensor Observation Services return sen-sor data as an O&M measurement. However, others use the 'SWE:dataArray' type. A response document with O&M measurements contains elements for individual observations. Every observation has an element for result, result time, UOM, procedure, feature of interest and observed property.

An SWE data array is an array of observations that share the same meta-data. For all observations that have the same feature of interest, procedure, observed property and UOM the result data is an array of results values

together with result or phenomenon times. The result value is separated from the result or phenomenon time using a predefined 'tokenseparator'. The combinations of result values and result or phenomenon times are separated using predefined 'blockseparators'.

The data from the received XML documents is directly added to an individual comma separated value string per combination of observed property and UOM. In case the temporal filter could not be used all data is looped over to remove observations outside the temporal range.

5.2.4 Data aggregation

After all observation data is retrieved it is first aggregated temporally. An empty dictionary is created to store each temporal granularity range that is inside the requested temporal range for all observations. A loop goes over the comma separated values and sorts them based on their result time per sensor location. The start time is subtracted from the result time, which results the time range from the start of the temporal range to the time of the observation. From this time range a modulo operation calculates how many times the temporal granularity fits in the time range between the start of the temporal range and the time of the observation. The start time is added to the temporal granularity times the outcome of the modulo operation to calculate the dictionary key to sort the observation by.

As soon as all the observations have been sorted the data is aggregated. For all values per key the average, minimum, maximum, median or sum is calculated. The resulting value replaces the values in the dictionary.

If spatial aggregation is part of the sensor data request this is performed after the temporal aggregation. Using Shapely's 9-intersection model functions the sensor locations are ordered per spatial feature. Finally, all values are aggregated per feature per temporal range.

5.3 SETTING UP THE WEB PROCESSING SERVICES

Creating two Web Processing Services using PyWPS.

Ouput XML using O&M schema or using JavaScript Object Notation (JSON) with the potential OGC canditate standard for JSON [Cox and Taylor, 2015].

Setting it up on the university server

5.4 CREATING A WEB APPLICATION FOR RETRIEVING SENSOR DATA

A web application has been created using the Web Processing Services from Paragraph 5.3.

6 | RESULTS

6.1 IMPLEMENTATION DIFFERENCES BETWEEN SENSOR OBSERVATION SERVICES

Even though both Sensor Observation Services use the SWE standards they are not exactly the same. They have different approaches for making identifiers and have slightly different content in their capabilities and describe sensor documents. This has caused the implemented to be more complex than expected. Especially providing a describe sensor document with as much information as possible (containing related features of interest, observed property and offerings) is important for making sense of the metadata. This should also be done using the appropriate SensorML tags, instead of the more general 'swe:keywords' for example.

When retrieving sensor data from the SOS using `GetObservation` requests one implementation provides this data as an array of comma separated values, using the SWE Array Observation class. The other implementation provides the same kind of data embedded in XML tags using the O&M Measurement class.

One of the SOS implementations allowed multiple observed properties to be related to an observation procedure. This was an unexpected finding, because a procedure is defined as a "method, algorithm or instrument, or system of these which may be used in making an observation" [ISO, 2011, p. 4]. When a procedure has different observed properties it is rather a combination of procedures. It is unlikely that the same procedure can observe one property at a certain time and a completely other property at another time. The original implementation of the WPS had to be adjusted to deal with this.

6.2 SEMANTICS IN SENSOR OBSERVATION SERVICES

6.2.1 mapping of observable properties

The metadata in Sensor Observation Services does not necessarily contain semantics. This is an issue when automatically retrieving metadata from different services. In the implementation this has caused a problem with identifying which observable properties in a SOS are the same as observable properties found in another SOS.

6.2.2 Automatically creating an URI scheme

If an URI is automatically created for data from an unknown source there is an amount of uncertainty to what the URI will look like. Either a random identifier can be assigned to it, or the identifier that is already provided by

the data source. This identifier most likely contains some reference to the nature of the real world thing the URI represents and is therefore preferable to a a completely random identifier. Adding the given identifier to the URI can create very long and strange URIs because different data sources have a different way of creating URIs since the O&M schema allows simply 'any URI'.

6.3 SPATIAL QUERIES WITH SPARQL

For retrieving data about a vector feature three methods for spatial querying have been implemented. First of all, spatial queries in which the server side receives the complete geometry of a feature and checks which point geometries it contains. Second of all, spatial queries in which the server side receives the bounding box of the complete geometry and check which points are within it. Third of all, spatial queries in which the server side receives the EEA reference grid cell that overlap the geometry and checks which points each cell contains. With the first method the client does not have to perform any spatial queries anymore. The second and third method will also return points that are not inside the geometry of the vector feature. These points could be filtered out by the client.

The Strabon and Parliament endpoint have been tested since they both handle GeoSPARQL queries. Strabon has been used in the final design, because the Parliament endpoint rejected certain longer queries (see 6.3).

Pubby software in combination with Apache Tomcat allows for a user interface that is easier to navigate through for humans. The links stored in RDF triples are represented as hyperlinks which can be used to navigate between pages about different concepts.

6.3.1 Vector queries

GeoSPARQL allows spatial queries in which a geometry can be inserted in the query. The implementation uses this functionality to test for spatial relations between geometries. For example between the point geometry of a sensor and the polygon geometry of an administrative unit. However, when geometries become more complicated their WKT definition becomes more verbose. This leads to the query being rejected based on its number of characters by the endpoint. This indicates that vector queries with complex geometries are not very efficient to include in a SPARQL query.

6.3.2 Raster queries

All raster cells are retrieved that overlap with the vector geometry. For these raster cell sensors are retrieved and later the excess ones are filtered out.

6.3.3 Bounding box queries

Creating a bounding box around a vector feature and make `GetObservation` to the SOS.

6.3.4   Latitude and longitude order

During the implementation a problem has come up regarding the order in which latitude and longitude are being presented in the SOS. The SOS of the RIVM provides point geometries in WGS84 as longitude, latitude and height. However, the Strabon endpoint expects the order to be latitude, longitude and height. This results false outcomes of spatial queries.

6.4   OUTPUT DATA

Data is outputted in XML or JSON according to the O&M schema. However, this schema cannot be used when different processes have been used by the data that has been aggregated.

6.5   COMPARING THE SENSOR INSTANCE REGISTRY WITH A SEMANTIC KNOWLEDGE BASE

Comparing Jirka and Nüst [2010] and Jirka and Bröring [2009] with the outcomes of Chapter 5

6.6   COMPARING THE SEMANTIC SENSOR MIDDLE-WARE WITH A SEMANTIC KNOWLEDGE BASE

Comparing the outcomes of Chapter 5 with methods for retrieving RDF from a SOS

7 | CONCLUSIONS

This thesis aimed to design a method that uses the semantic web to improve sensor data discovery as well as the integration and aggregation of sensor data from multiple sources. Each of the four subquestions posed in the introduction will be answered in this chapter before addressing the main research question.

SUBQUESTION 1. TO WHAT EXTENT CAN SENSOR METADATA BE AUTOMATICALLY RETRIEVED FROM ANY SOS?

Data inside a Sensor Observation Service can be automatically retrieved from a SOS, using the methods described in Paragraph 4.1.1. However, there are two things in the SWE standards that should be improved to make the proposed design work better.

First of all, the capabilities document is currently only required to contain a list of all features of interest as a parameter for the `GetObservation` request. Optionally the features of interest can also be mentioned as metadata per offering. In the case of air quality these features represent the sensor locations. However, it is merely required to list the URIs of the features of interest. To retrieve their geometries either `GetFeatureOfInterest` or `GetObservation` requests have to be made. This is especially cumbersome since the `GetFeatureOfInterest` response does not link features of interest to procedures or observed properties. This relation is therefore only visible by either combining the observed properties and procedures per offering to the `GetFeatureOfInterest` response or by requesting observations from each sensor location. Therefore, I propose that the capabilities document should not only list the URIs of the features of interest, but also describe the geometry and observed properties of each of them.

Secondly, in the current implementation of the XML schemas for SOS (http://schemas.opengis.net/sos/2.0) and O&M (http://schemas.opengis.net/om/2.0/) identifiers for features of interest, observed properties and procedures can be 'any URI'. This means that an URL with semantics can be provided, but a non-semantic URI would also be valid. I propose to make the standard more strict and require a semantic URL that can be resolved to an RDF document. Without these semantics it is hard to use a SOS for both humans and automatic processes, especially when two or more are used in combination with each other. Paragraph 4.1.3 describes how to use a PURL server for creating persistent semantic URLs that resolve to an RDF document.

Besides the two changes that should be made to the SWE standards there is another issue with automatically retrieving sensor metadata from a SOS, namely the order of coordinates for a feature of interest. Both Sensor Observation Services used in this thesis had a different order for the latitude

51

and the longitude of their point coordinates. This is a standardisation issue already identified by many authors, which should be decided upon by the geomatics and geoscience community. For the method described in this thesis it is irrelevant which order is being used, as long as its clearly described in individual cases or prescribed using an international standard.

SUBQUESTION 2. TO WHAT EXTENT CAN SENSOR METADATA FROM A SOS BE AUTOMATICALLY CONVERTED TO LINKED DATA AND PUBLISHED ON THE SEMANTIC WEB?

In a SOS there are XML schemas that contain general semantics about the metadata. They identify what the different URIs represent (e.g. observed properties, procedures, features of interest). The om-lite and sam-lite ontologies have been used to make linked data from this metadata. However, a number of classes should be added to these ontologies to make it suit this process better. First of all, a class is required that distinguishes the *process* of creating an observation from the physical *device* that uses this process. This 'sensor' class could be modelled as a device that uses a procedure at a certain sampling point. Adding this class takes away some of the ambiguity between defined processes and actually deployed sensors. Therefore, it will be easier to perform (spatial) SPARQL queries that return deployed sensors of which data can be retrieved.

Another class that should be represented in an ontology is the Sensor Observation Service. The current prototype design (Chapter 5) used a single endpoint for storing all metadata from Sensor Observation Services. Therefore, it was known that the source of data is a SOS, but it has not been properly defined in a linked data ontology. If programs crawling the semantic web can identify a data source such as a SOS and understand its allowed queries (e.g. `GetCapabilities`, `DescribeSensor`, `GetObservation`), they can retrieve data from it without requiring prior knowledge. This way sensor data using other platforms such as the SensorThings API could be discovered and retrieved in the same way and used in combination with each other. Therefore the extensions of current ontologies is further discussed as future research in Chapter 9.

SUBQUESTION 3. WHAT IS AN EFFECTIVE BALANCE BETWEEN THE SEMANTIC WEB AND THE GEO WEB IN THE CHAIN OF DISCOVERING, RETRIEVING AND PROCESSING SENSOR DATA?

A number of authors have shown that triple stores do not perform as good as databases when their data is requested via a web application. On the other hand, linked data is very well suited for discovering data as it is literally 'linked' to related data. Therefore, this thesis aimed to design a method for using the semantic web in combination with sensor web applications, where the semantic web contains metadata and the geoweb observation data.

However, there is a grey area of functions that could be implemented using either on of these two parts of the web.

For example, the semantic web could have a bounding box per SOS containing all features of interest it offers in combination with a list of all observed properties. In this case spatial and temporal filters would have to be applied at the SOS side when retrieving observation data. On the other hand, the semantic web could also contain detailed information about individual sensors. This way the Sensor Observation Services are only used to retrieve observation data of already selected sensors.

Both options have been considered. The second option has been used in this thesis for a number of reasons. First of all, it was found that not all Sensor Observation Services offer the same filter capabilities. In order for the first option to be viable every SOS should have a minimum amount of filter capabilities implemented by default. Second of all, if a user is interested in sensors located in a specific area the bounding box of all sensors might be misleading. If there is an outlying sensor or if all sensors are inside a curved or diagonal vector geometry a SOS might seem relevant while it actually is not. The result of this is that many unnecessary requests will be send to a SOS keeping the user waiting and the SOS server busy for no reason.

Third of all, semantic information about specific sensors can be linked to by other related linked data. This can be done by the organisation maintaining the sensor or by other organisations. For example links to the manufacturer, the quality of observations achieved by a certain model of sensors, or the conditions under which the sensor is placed could be useful for anyone interested in the observation data.

However, the down side of the second approach is that the performance is lower. The SPARQL query language is not as fast as for example queries using PostgreSQL and Postgis. Adding spatial filters is possible, but is not very fast for most vector geometries. Furthermore, a larger amount of data needs to be send over the web in the process of discovering sensors. Since GeoSPARQL and stSPARQL uses the verbose WKT or GML encodings the queries can also be rejected for exceeding the maximum amount of characters. For this reason it is found that spatial filters using vector geometries should best be avoided.

Still, discovering sensors is only a matter of seconds using the second approach. Automatically retrieving sensor data can take up to a couple minutes depending on the amount of sensors for which data is requested and the temporal range. However, it should be noted that performance optimization has not been a part of this thesis. It is likely that this can still be improved significantly (see Chapter 9).

SUBQUESTION 4. TO WHAT EXTENT CAN ALREADY EXISTING STANDARDS FOR RETRIEVING DATA BE (RE)USED FOR A SERVICE THAT SUPPLIES INTE- GRATED AND AGGREGATED SENSOR DATA?

All data models and services in this thesis have been used because they are based on open standards. Designs for two processes have been explored: an automated process for creating linked data from metadata in a SOS and a process for discovering, retrieving and processing sensor data. These pro-

cesses were created using OGC Web Processing Services, which is a standard API for spatial data processes on the web. The WPS is well suited for these two applications.

The O&M standard could be reused on the semantic web using the om-lite and sam-lite ontologies. These are lightweight linked data ontologies based on O&M. Performing spatial queries on this linked data is possible using OGC's GeoSPARQL as well as using Strabon's stRDF (Paragraph 3.3).

The O&M observation schema is being reused for the spatially aggregated sensor data that is retrieved from the discovered Sensor Observation Services. However, the schema only allows sensor data from the same procedure. The result of this is that only observations from the same SOS fit in the O&M observation schema, as different procedures were implemented by the different organisations maintaining a SOS. Therefore, I propose to either allow the procedure to be nullable in the schema or to allow the procedure element to contain an array of multiple procedures.

MAIN QUESTION: TO WHAT EXTENT CAN THE SEMANTIC WEB IMPROVE THE DISCOVERY, INTEGRATION AND AGGREGATION OF DISTRIBUTED SENSOR DATA?

In this thesis a method has been designed to create an online knowledge base with linked metadata extracted from Sensor Observation Services. This helps discovering, integrating and aggregation sensor data, while for efficient data retrieval the SOS is still used. The results show that such a knowledge base makes it easy to discover sensors and their corresponding sources. It allows for online processes to automatically retrieve and process sensor data. It can be generated from any SOS, if there are a minimum of semantics provided. Most SWE standards can be (re)used, although some need to be changed for the method presented here to work for every implementation of a SOS. Observed properties and procedures need to be semantically defined by the organisation maintaining a SOS and linked data ontologies should be extended to define the Sensor Observation Service and its offerings.

A number of design decisions have been made which will be further discussed in Chapter 8. Areas for improving and extending the presented methods are described in Chapter 9.

8 | DISCUSSION

8.1 METADATA DUPLICATION

The method presented in this thesis takes metadata from a SOS, converts it to linked data and publishes it on the semantic web. Although the metadata has taken another form, it is now stored twice in two different locations. This may not be desirable, for instance when data is updated in the original source and its linked data equivalent is not. Also more storage space is required for the same amount of data. However, extra functionality is achieved in return.

8.2 METADATA QUALITY

The quality of the metadata in the SOS influences the quality of the metadata SPARQL endpoint.

8.3 AUTOMATED PROCESS

If there is no meaning added to definitions like observed property, the metadata is not machine understandable. In this case, manual work has to be done to make it machine understandable. Only after this manual process it can be published on the semantic web.

8.4 EXPLICIT TOPOLOGICAL RELATIONS

Spatial features have topological relations with other spatial features. These relations can be made explicit on the semantic web. However, in this thesis they have not been made explicit and are calculated on-the-fly with spatial queries using GeoSPARQL. Making topological relations explicit in a subject-predicate-object structure could improve query speed, as they are likely less expensive than spatial queries. However, this is a trade-off with the required storage space. Furthermore, the chances of incorrect or broken links increase as both features and topological relations can change over time.

8.5 THE USE OF A CATALOG SERVICE

The methods presented in Chapter 5 could also be implemented on top of a CSW: the metadata from a SOS would be inside a CSW and this CSW would have a SPARQL endpoint connected to it. Describe pro's and con's.

9 FUTURE RESEARCH

Include more sensor standards besides SOS, such as tinySOS and Sensor-Things API.

Include more geoweb standards such as WFS, WCS, etc.

Extend current ontologies with semantics about OGC geo web services such as SOS and their allowed requests.

Improving performance: currently discovering sensors is fast for bounding box or raster queries. Vector queries take up more time, but is still relatively fast. However, real performance improvements could still be gained in the automatic querying of a SOS. Currently for every sensor a request is send, but this could probably be improved by requesting all sensor data from the same SOS using fewer or perhaps even using only a single request.

A | DATA VISUALISATIONS

**Figure A.1:** Dataset of municipalities in the Netherlands and Belgium in 2015 (from Dutch cadaster and GADM.org)

Figure A.2: Dataset of provinces in the Netherlands and Belgium in 2015 (from Dutch cadaster and GADM.org)

Figure A.3: Dataset of the Netherlands and Belgium in 2015 (from GADM.org)
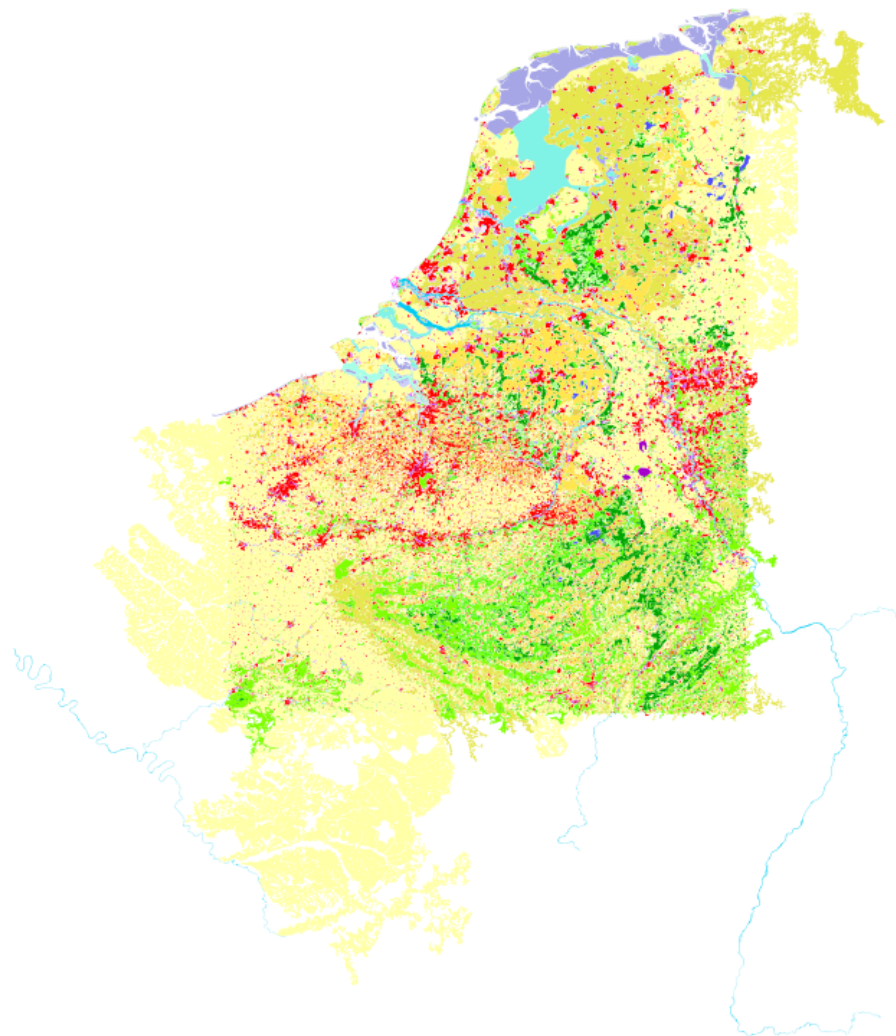
Figure A.4: Dataset of landcover in the Netherlands and Belgium in 2012 (from Copernicus The European Earth Observation Programme)

CORINE 2012 dataset of landcover in
South Holland

Legend

CORINE 2012

- Continuous urban fabric
- Discontinuous urban fabric
- Industrial or commercial units
- Road and rail networks and associated land
- Port areas
- Airports
- Mineral extraction sites
- Dump sites
- Construction sites
- Green urban areas
- Sport and leisure facilities
- Non-irrigated arable land
- Permanently irrigated land
- Rice fields
- Vineyards
- Fruit trees and berry plantations
- Olive groves
- Pastures
- Annual crops associated with permanent crops
- Complex cultivation patterns
- Land principally occupied by agriculture with significant areas of natural vegetation
- Agro-forestry areas
- Broad-leaved forest
- Coniferous forest
- Mixed forest
- Natural grasslands
- Moors and heathland
- Sclerophyllous vegetation
- Transitional woodland-shrub
- Beaches - dunes - sands
- Bare rocks
- Sparsely vegetated areas
- Burnt areas
- Glaciers and perpetual snow
- Inland marshes
- Peat bogs
- Salt marshes
- Salines
- Intertidal flats
- Water courses
- Water bodies
- Coastal lagoons
- Estuaries
- Sea and ocean
- NODATA

Figure A.5: Landcover of the province of South Holland (subsection of the dataset from Figure A.4)

Figure A.6: EEA reference grid cells with a resolution of 100km² overlapping the Netherlands and Belgium

Figure A.7: EEA reference grid cells with a resolution of 10km² overlapping the Netherlands and Belgium

Figure A.8: Webmap by the RIVM showing their air quality sensor network (http://www.lml.rivm.nl/meetnet)

Figure A.9: Webmap by IRCEL-CELINE showing their air quality sensor network (http://www.irceline.be/en/air-quality/measurements/monitoring-stations/)

Figure A.10: Google Streetview image of RIVM sensor location in Amsterdam in 2015

B | WEB PROCESSING SERVICE RESPONSE DOCUMENTS

B.1 EXAMPLE CAPABILITIES DOCUMENT

```
<wps:Capabilities xmlns:xlink="http://www.w3.org/1999/xlink"
   xmlns:wps="http://www.opengis.net/wps/1.0.0"
   xmlns:ows="http://www.opengis.net/ows/1.1"
   xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
   service="WPS" version="1.0.0" xml:lang="en-CA"
   xsi:schemaLocation="http://www.opengis.net/wps/1.0.0
   http://schemas.opengis.net/wps/1.0.0/wpsGetCapabilities_response.xsd"
   updateSequence="1">
 <ows:ServiceIdentification>
    <ows:Title>PyWPS Server</ows:Title>
    <ows:Abstract>
       See http://pywps.wald.intevation.org and
          http://www.opengeospatial.org/standards/wps
    </ows:Abstract>
    <ows:Keywords>
       <ows:Keyword>WPS</ows:Keyword>
       <ows:Keyword>SWE</ows:Keyword>
       <ows:Keyword>SOS</ows:Keyword>
    </ows:Keywords>
    <ows:ServiceType>WPS</ows:ServiceType>
    <ows:ServiceTypeVersion>1.0.0</ows:ServiceTypeVersion>
    <ows:Fees>None</ows:Fees>
    <ows:AccessConstraints>none</ows:AccessConstraints>
 </ows:ServiceIdentification>
 <ows:ServiceProvider>
    <ows:ProviderName>Delft University of
       Technology</ows:ProviderName>
    <ows:ProviderSite
       xlink:href="http://masterthesistudelft.herokuapp.com/"/>
    <ows:ServiceContact>
    <ows:IndividualName>Ivo de Liefde</ows:IndividualName>
    <ows:PositionName>MSc. student Geomatics for the Built
       Environment</ows:PositionName>
    <ows:ContactInfo>
      <ows:Address>
      <ows:DeliveryPoint>Julianalaan 134</ows:DeliveryPoint>
      <ows:City>Delft</ows:City>
      <ows:PostalCode>2628 BL</ows:PostalCode>
      <ows:Country>the Netherlands</ows:Country>
      <ows:ElectronicMailAddress>i.deliefde@student.tudelft.nl</ows:ElectronicMailAddress>
      </ows:Address>
```

```xml
            <ows:OnlineResource
                xlink:href="http://masterthesistudelft.herokuapp.com/"/>
            <ows:HoursOfService>0:00-24:00</ows:HoursOfService>
            <ows:ContactInstructions>none</ows:ContactInstructions>
        </ows:ContactInfo>
        <ows:Role>
            Created WPS
        </ows:Role>
        </ows:ServiceContact>
    </ows:ServiceProvider>
    <ows:OperationsMetadata>
        <ows:Operation name="GetCapabilities">
            <ows:DCP>
                <ows:HTTP>
                    <ows:Get
                        xlink:href="http://localhost/cgi-bin/wps?"/>
                    <ows:Post
                        xlink:href="http://localhost/cgi-bin/wps"/>
                </ows:HTTP>
            </ows:DCP>
        </ows:Operation>
        <ows:Operation name="DescribeProcess">
            <ows:DCP>
                <ows:HTTP>
                    <ows:Get
                        xlink:href="http://localhost/cgi-bin/wps?"/>
                    <ows:Post
                        xlink:href="http://localhost/cgi-bin/wps"/>
                </ows:HTTP>
            </ows:DCP>
        </ows:Operation>
        <ows:Operation name="Execute">
            <ows:DCP>
                <ows:HTTP>
                    <ows:Get
                        xlink:href="http://localhost/cgi-bin/wps?"/>
                    <ows:Post
                        xlink:href="http://localhost/cgi-bin/wps"/>
                </ows:HTTP>
            </ows:DCP>
        </ows:Operation>
    </ows:OperationsMetadata>
    <wps:ProcessOfferings>
        <wps:Process wps:processVersion="1.0">
            <ows:Identifier>LinkedDataFromSOS</ows:Identifier>
            <ows:Title>Creates Linked Data of SOS
                metadata</ows:Title>
            <ows:Abstract>
                This process takes an HTTP address of a Sensor
                    Observation Service (SOS) as input and converts
                    the metadata to linked data.
            </ows:Abstract>
        </wps:Process>
```

```xml
    <wps:Process wps:processVersion="1.0">
       <ows:Identifier>GetSensors</ows:Identifier>
       <ows:Title>
          Automatically retrieves sensors from heterogenous
             sources using the semantic web
       </ows:Title>
       <ows:Abstract>
          This process takes a sensor data request with
             parameters for spatial features of interest,
             observed property, temporal range and
             granularity, and finds all relevant sensor data
             sources on the semantic web.
       </ows:Abstract>
    </wps:Process>
    <wps:Process wps:processVersion="1.0">
       <ows:Identifier>GetSensorData</ows:Identifier>
       <ows:Title>
          Automatically retrieves, integrates and aggregates
             heterogenous sensor data using the semantic web
       </ows:Title>
       <ows:Abstract>
       This process takes sensors found by the WPS
          'GetSensors' and automatically integrates and
          aggregates the data from different sources on the
          web.
       </ows:Abstract>
    </wps:Process>
  </wps:ProcessOfferings>
  <wps:Languages>
    <wps:Default>
       <ows:Language>en-CA</ows:Language>
    </wps:Default>
    <wps:Supported>
       <ows:Language>en-CA</ows:Language>
    </wps:Supported>
  </wps:Languages>
  <wps:WSDL xlink:href="http://localhost/cgi-bin/wps?WSDL"/>
</wps:Capabilities>
```

B.2 EXAMPLE DESCRIBE PROCESS DOCUMENT

```
<wps:ProcessDescriptions
    xmlns:wps="http://www.opengis.net/wps/1.0.0"
    xmlns:ows="http://www.opengis.net/ows/1.1"
    xmlns:xlink="http://www.w3.org/1999/xlink"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xsi:schemaLocation="http://www.opengis.net/wps/1.0.0
    http://schemas.opengis.net/wps/1.0.0/wpsDescribeProcess_response.xsd"
    service="WPS" version="1.0.0" xml:lang="en-CA">
<ProcessDescription wps:processVersion="1.0"
    storeSupported="true" statusSupported="false">
  <ows:Identifier>LinkedDataFromSOS</ows:Identifier>
  <ows:Title>Creates Linked Data of SOS metadata</ows:Title>
  <ows:Abstract>
     This process takes an HTTP address of a Sensor
         Observation Service (SOS) as input and converts the
         metadata to linked data.
  </ows:Abstract>
  <DataInputs>
     <Input minOccurs="0" maxOccurs="1">
        <ows:Identifier>observed_properties</ows:Identifier>
        <ows:Title>
           Input link to turtle file with mappings of
               observed property identifiers to DBPedia URIs
        </ows:Title>
        <LiteralData>
           <ows:DataType
               ows:reference="http://www.w3.org/TR/xmlschema-2/#string">string</
           <ows:AnyValue/>
           <DefaultValue>http://inspire.rivm.nl/sos/eaq/service?</DefaultValue>
        </LiteralData>
     </Input>
     <Input minOccurs="0" maxOccurs="1">
        <ows:Identifier>input_url</ows:Identifier>
        <ows:Title>
           Input a string containing an HTTP address of a
               Sensor Observation Service (SOS). For
               example: 'http://someaddress.com/sos?'
        </ows:Title>
        <LiteralData>
           <ows:DataType
               ows:reference="http://www.w3.org/TR/xmlschema-2/#string">string</
           <ows:AnyValue/>
           <DefaultValue>http://inspire.rivm.nl/sos/eaq/service?</DefaultValue>
        </LiteralData>
     </Input>
  </DataInputs>
</ProcessDescription>
</wps:ProcessDescriptions>
```

B.3 EXAMPLE EXECUTE DOCUMENT

```xml
<wps:ExecuteResponse
    xmlns:wps="http://www.opengis.net/wps/1.0.0"
    xmlns:ows="http://www.opengis.net/ows/1.1"
    xmlns:xlink="http://www.w3.org/1999/xlink"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xsi:schemaLocation="http://www.opengis.net/wps/1.0.0
    http://schemas.opengis.net/wps/1.0.0/wpsExecute_response.xsd"
    service="WPS" version="1.0.0" xml:lang="en-CA"
    serviceInstance="http://localhost/cgi-bin/wps?service=WPS&request=GetCapabilities&version=1.
    statusLocation="http://localhost/wps/wpsoutputs/pywps-146072575066.xml">
  <wps:Process wps:processVersion="1.0">
    <ows:Identifier>GetSensors</ows:Identifier>
    <ows:Title>
       Automatically retrieves sensors from heterogenous
           sources using the semantic web
    </ows:Title>
    <ows:Abstract>
       This process takes a sensor data request with
           parameters for spatial features of interest,
           observed property, temporal range and granularity,
           and finds all relevant sensor data sources on the
           semantic web.
    </ows:Abstract>
  </wps:Process>
  <wps:Status creationTime="2016-04-15T15:09:52Z">
    <wps:ProcessSucceeded>PyWPS Process GetSensors
       successfully calculated</wps:ProcessSucceeded>
  </wps:Status>
  <wps:ProcessOutputs>
    <wps:Output>
       <ows:Identifier>output</ows:Identifier>
       <ows:Title>Output sensor data</ows:Title>
       <wps:Data>
          <wps:ComplexData mimeType="text/JSON">
             OUTPUT JSON DATA
          </wps:ComplexData>
       </wps:Data>
    </wps:Output>
  </wps:ProcessOutputs>
</wps:ExecuteResponse>
```

BIBLIOGRAPHY

52 North (2016). Sensor discovery. [online] http://52north.org/communities/sensorweb/discovery [accessed on April, 20th, 2016].

Atkinson, R. A., Taylor, P., Squire, G., Car, N. J., Smith, D., and Menzel, M. (2015). Joining the Dots: Using Linked Data to Navigate between Features and Observational Data. In *Environmental Software Systems. Infrastructures, Services and Applications*, pages 121–130. Springer.

Atzori, L., Iera, A., and Morabito, G. (2010). The internet of things: A survey. *Computer networks*, 54(15):2787–2805.

Barnaghi, Payam and Wang, Wei and Henson, Cory and Taylor, Kerry (2012). Semantics for the Internet of Things: early progress and back to the future. *International Journal on Semantic Web and Information Systems (IJSWIS)*, 8(1):1–21.

Battle, R. and Kolas, D. (2012). Enabling the geospatial semantic web with parliament and geosparql. *Semantic Web*, 3(4):355–370.

Beckett, D., Berners-Lee, T., Prud'hommeaux, E., and Carothers, G. (2014). W3C RDF 1.1 Turtle. [online] http://www.w3.org/TR/turtle/ [accessed on December 9th, 2015].

Berners-Lee, T. and Connolly, D. (2011). W3C Notation3 (N3): A readable RDF syntax. [online] http://www.w3.org/TeamSubmission/n3/ [accessed on December 9th, 2015].

Berners-Lee, T., Hendler, J., Lassila, O., et al. (2001). The semantic web. *Scientific american*, 284(5):28–37.

Bizer, C., Heath, T., and Berners-Lee, T. (2009). Linked data-the story so far. *Semantic Services, Interoperability and Web Applications: Emerging Concepts*, pages 205–227.

Botts, M., Percivall, G., Reed, C., and Davidson, J. (2007). OGC Sensor Web Enablement: Overview And High Level Architecture. OGC document 06-021r1.

Botts, M., Percivall, G., Reed, C., and Davidson, J. (2008). OGC sensor web enablement: Overview and high level architecture. In *GeoSensor networks*, pages 175–190. Springer.

Bröring, A., Stasch, C., and Echterhoff, J. (2012). OGC Sensor observation service interface standard.

Cambridge Semantics (2015). Introduction to the Semantic Web. [online] https://www.cambridgesemantics.com/semantic-university/introduction-semantic-web [accessed on December 8th, 2015].

Compton, M., Barnaghi, P., Bermudez, L., GarcíA-Castro, R., Corcho, O., Cox, S., Graybeal, J., Hauswirth, M., Henson, C., Herzog, A., et al.

(2012). The SSN ontology of the W3C semantic sensor network incubator group. *Web Semantics: Science, Services and Agents on the World Wide Web*, 17:25–32.

Corcho, O. and Garcia-Castro, R. (2010). Five challenges for the Semantic Sensor Web. *Semantic Web-Interoperability, Usability, Applicability*, 1.1(2):121–125.

Cox, S. J. D. (2015a). Observations and Sampling. [online] `https://www.seegrid.csiro.au/wiki/AppSchemas/ObservationsAndSampling` [accessed on December 1st, 2015].

Cox, S. J. D. (2015b). Ontology for observations and sampling features, with alignments to existing models.

Cox, S. J. D. (2015c). OWL for Observations. [online] `http://def.seegrid.csiro.au/ontology/om/om-lite` [accessed on November 24th, 2015].

Cox, S. J. D. (2015d). OWL for Sampling Features. [online] `http://def.seegrid.csiro.au/ontology/om/sam-lite` [accessed on November 24th, 2015].

Cox, S. J. D. and Taylor, P. (2015). OGC Observations and Measurements — JSON implementation.

Cyganiak, R., Wood, D., and Lanthaler, M. (2014). RDF 1.1 Concepts and Abstract Syntax. [online] `https://www.w3.org/TR/rdf11-concepts` [accessed on February 2nd, 2016].

Gandon, F. and Schreiber, G. (2014). W3C RDF 1.1 XML Syntax. [online] `http://www.w3.org/TR/rdf-syntax-grammar/` [accessed on December 9th, 2015].

Ganesan, D., Ratnasamy, S., Wang, H., and Estrin, D. (2004). Coping with irregular spatio-temporal sampling in sensor networks. *ACM SIGCOMM Computer Communication Review*, 34(1):125–130.

Henson, C., Pschorr, J. K., Sheth, A. P., Thirunarayan, K., et al. (2009). SemSOS: Semantic sensor observation service. In *Collaborative Technologies and Systems, 2009. CTS'09. International Symposium on*, pages 44–53. IEEE.

Hu, C., Guan, Q., Chen, N., Li, J., Zhong, X., and Han, Y. (2014). An Observation Capability Metadata Model for EO Sensor Discovery in Sensor Web Enablement Environments. *Remote Sensing*, 6(11):10546–10570.

INSPIRE (2014). Guidelines for the use of Observations & Measurements and Sensor Web Enablement-related standards in INSPIRE Annex II and III data specification development.

INSPIRE (2015). INSPIRE Roadmap. [online] `http://inspire.ec.europa.eu/index.cfm/pageid/44` [accessed on December 2nd, 2015].

International Organisation for Standardisation (2005). ISO 19109:2005; Geographic information — Rules for application schema.

International Organisation for Standardisation (2007). ISO 19136:2007; Geographic information — Geography Markup Language (GML).

ISO (2011). ISO 19156:2011; Geographic information – Observations and measurements. [online] `http://www.iso.org/iso/iso_catalogue/catalogue_tc/catalogue_detail.htm?csnumber=32574` [accessed on December 2 nd, 2015].

Janowicz, K., Broring, A., Stasch, C., Schad, S., Everding, T., and Llaves, A. (2013). A RESTful Proxy and Data Model for Linked Sensor Data. *International Journal of Digital Earth*, 6(3):233–254.

Jazayeri, Mohammad Ali and Liang, Steve HL and Huang, Chih-Yuan (2015). Implementation and Evaluation of Four Interoperable Open Standards for the Internet of Things. *Sensors*, 15(9):24343–24373.

Ji, C., Liu, J., and Wang, X. (2014). A Review for Semantic Sensor Web Research and Applications. *Advanced Science and Technology Letters*, 48:31–36.

Jirka, S. and Bröring, A. (2009). OGC Sensor Observable Registry Discussion Paper. Reference number: OGC 09-112.

Jirka, S. and Nüst, D. (2010). OGC Sensor Instance Registry Discussion Paper. Reference number: OGC 10-171.

Koivunen, M.-R. and Miller, E. (2002). W3C Semantic Web Activity. In Hyvönen, E., editor, *Semantic Web Kick-off Seminar in Finland*.

Korteweg, P., Marchetti-Spaccamela, A., Stougie, L., and Vitaletti, A. (2007). *Data aggregation in sensor networks: Balancing communication and delay costs*. Springer.

Koubarakis, M. and Kyzirakos, K. (2010). Modeling and querying metadata in the semantic sensor web: The model stRDF and the query language stSPARQL. In *The semantic web: research and applications*, pages 425–439. Springer.

Lassila, O. and Swick, R. R. (1999). Resource Description Framework (RDF) Model and Syntax Specification. [online] `http://www.w3.org/TR/PR-rdf-syntax/` [accessed on December 8th, 2015].

Lebo, T., Sahoo, S., and McGuinness, D. (2013). PROV-O: The PROV Ontology. [online] `http://www.w3.org/TR/prov-o/` [accessed on December 11th, 2015].

Manola, F., Miller, E., and McBride, B. (2014). W3C RDF Primer. [online] `http://www.w3.org/TR/rdf11-primer/` [accessed on December 9th, 2015].

Missier, G. A. (2015). Towards a Web application for viewing Spatial Linked Open Data of Rotterdam. Master's thesis, Delft University of Technology.

Moir, E., Moonen, T., and Clark, G. (2014). What are Future Cities: Origins, Meanings and Uses.

Nebert, D., Whiteside, A., and Vretanos, P. (2007). Opengis catalogue services specification.

OGC (2014). OGC SensorML: Model and XML Encoding Standard.

Open Geospatial Consortium (2007). OGC Catalogue Services Specification.

Open Geospatial Consortium (2015). OGC WPS 2.0 Interface Standard.

OWL working group (2012). Web Ontology Language (OWL). [online] http://www.w3.org/2001/sw/wiki/OWL [accessed on December 18[th], 2015].

Percivall, G. (2015). OGC Smart Cities Spatial Information Framework. OGC Internal reference number: 14-115.

Perry, M. and Herring, J. (2012). GeoSPARQL - A Geographic Query Language for RDF Data.

Price Waterhouse Coopers (2014). Sensing the future of the Internet of Things. [online] https://www.pwc.com/us/en/increasing-it-effectiveness/assets/future-of-the-internet-of-things.pdf [accessed on December 18[th], 2015].

Pschorr, J., Henson, C. A., Patni, H. K., and Sheth, A. P. (2010). Sensor discovery on linked data.

Pschorr, J. K. (2013). SemSOS: an Architecture for Query, Insertion, and Discovery for Semantic Sensor Networks. Master's thesis, Wright State University.

PURL (2016). Batch Uploading to a PURL Server v1.0-1.6.x. [online] https://code.google.com/archive/p/persistenturls/wikis/PURLBatchUploadingVersionOne.wiki [accessed on February 19[th], 2016].

Shafer, K., Weibel, S., Jul, E., and Fausey, J. (2016). Introduction to Persistent Uniform Resource Locators. [online] https://purl.oclc.org/docs/long_intro.html [accessed on February 18[th], 2016].

Sheth, A., Henson, C., and Sahoo, S. S. (2008). Semantic Sensor Web. *IEEE Internet Computing*, 12(4):78–83.

Stasch, C., Autermann, C., Foerster, T., and Pebesma, E. (2011a). Towards a spatiotemporal aggregation service in the sensor web. Poster presentation. In *The 14th AGILE International Conference on Geographic Information Science*.

Stasch, C., Schade, S., Llaves, A., Janowicz, K., and Bröring, A. (2011b). Aggregating linked sensor data. In Taylor, K., Ayyagari, A., and de Roure, D., editors, *Proceedings of the 4th International Workshop on Semantic Sensor Networks*, page 46.

Stasch, C., Scheider, S., Pebesma, E., and Kuhn, W. (2014). Meaningful spatial prediction and aggregation. *Environmental Modelling & Software*, 51:149–165.

Strobl, C. (2008). *Dimensionally Extended Nine-Intersection Model (DE-9IM)*. Springer.

Theunisse, I. A. H. (2015). The Visualization of Urban Heat Island Indoor Temperatures. Master's thesis, TU Delft, Delft University of Technology.

van der Hoeven, F., Wandl, A., Demir, B., Dikmans, S., Hagoort, J., Moretto, M., Sefkatli, P., Snijder, F., Songsri, S., Stijger, P., et al. (2014). Sensing Hotterdam: Crowd sensing the Rotterdam urban heat island. *SPOOL*, 1(2):43–58.

Van der Hoeven, F. D. and Wandl, A. (2015). Hotterdam: How space is making Rotterdam warmer, how this affects the health of its inhabitants, and what can be done about it. Technical report, TU Delft, Faculty of Architecture and the Built Environment.

W3C Semantic Sensor Network Incubator Group (2011). Semantic Sensor Network Ontology. [online] `http://www.w3.org/2005/Incubator/ssn/ssnx/ssn` [accessed on December 9th, 2015].

Wang, M., Perera, C., Jayaraman, P. P., Zhang, M., Strazdins, P., and Ranjan, R. (2015a). City Data Fusion: Sensor Data Fusion in the Internet of Things.

Wang, X., Zhang, X., and Li, M. (2015b). A Review of Studies on Semantic Sensor Web. *Advanced Science and Technology Letters*, 83:94–97.

Xiang, L., Luo, J., and Rosenberg, C. (2013). Compressed data aggregation: Energy-efficient and high-fidelity data collection. *Networking, IEEE/ACM Transactions on*, 21(6):1722–1735.

Zanella, A., Bui, N., Castellani, A., Vangelista, L., and Zorzi, M. (2014). Internet of things for smart cities. *Internet of Things Journal, IEEE*, 1(1):22–32.